

432 Week 4 Slides

github.com/THOMASELOVE/2020-432

2020-02-04 & 02-06

This Week's Agenda

- County Health Rankings Data
- Pre-processing of data
- Investigating an outcome transformation
- Model (Variable) Selection
 - Why Stepwise Regression is Terrible
 - How to do “Best Subsets” Linear Regression
- Cross-Validation
 - Validation Set approach
 - K-fold Cross-Validation
- and perhaps some more. . .

Setup

```
library(here); library(magrittr); library(janitor)
library(patchwork); library(naniar)
library(knitr); library(mosaic); library(skimr)

library(GGally)
library(car)
library(leaps)
library(caret)
library(modelr)

library(broom)
library(tidyverse)

theme_set(theme_bw())
```

US County Health Rankings, 2017

US County Health Rankings Data

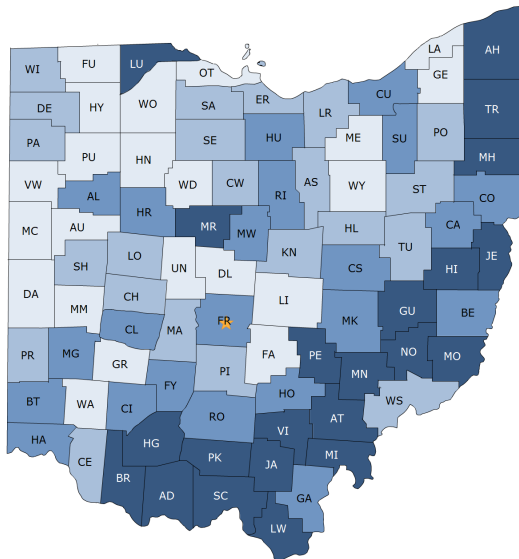
- We'll use the 2017 data in our examples, to leave more recent data available for you in Projects and future work.
- Data source is [this link](https://www.countyhealthrankings.org) at <https://www.countyhealthrankings.org>.

Next two slides show Ohio counties in 2017 and their rankings on:

- Health Factors: weighted scores for health behaviors, clinical care, social and economic factors, and the physical environment (Cuyahoga ranked 56th out of the 88 Ohio counties)
- Health Outcomes: equal weighting of length and quality of life (Cuyahoga ranked 65th)

In each slide, lighter shades indicate better performance.

2017 Health Factors - Ohio

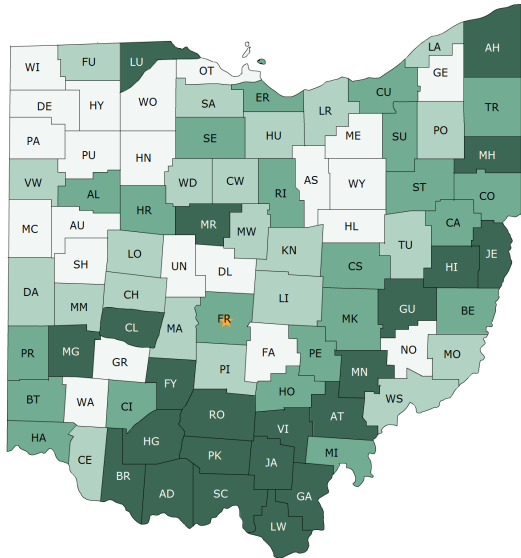


Rank 1-22 Rank 23-44 Rank 45-66 Rank 67-88

County Health
Rankings & Roadmaps
Building a Culture of Health, County by County

A Robert Wood Johnson Foundation program

2017 Health Outcomes - Ohio



Rank 1-22 Rank 23-44 Rank 45-66 Rank 67-88

**County Health
Rankings & Roadmaps**
Building a Culture of Health, County by County

A Robert Wood Johnson Foundation program

Data Ingest and Cleanup

```
count_2017 <-  
  read_csv(here("data/countyhealthrankings_2017.csv")) %>%  
  clean_names() %>%  
  type.convert() %>%  
  mutate(fips = as.character(fips),  
         state = as.character(state),  
         county = as.character(county))  
  
dim(count_2017)
```

```
[1] 3136 34
```


Codebook (2017 County Health Rankings), I

Variable	Description
fips	FIPS code (5 digits: 1-2: state, 3-5: county)
state	State Name
county	County Name
yrs_lost_rate	Years of potential life lost before age 75 per 100,000 population (age-adjusted, 2012-14)
population	County population, Census Population Estimates, 2015
not_english	% of county residents preferring language besides English (ACS 2010)
hh_income	Median Household Income in \$1000s (Small Area Income & Poverty Estimates 2015)
income_ratio	Ratio of household income at the 80th percentile to income at the 20th percentile (ACS 2011-15)
health_costs	Average per-capita health care costs (Dartmouth Atlas of Health Care 2014)

Codebook (2017 County Health Rankings), II

Variable	Description
social_assoc	Membership associations per 10,000 population (County Business Partners 2014)
pct_smokers	% of adults who currently smoke (2015 BRFSS)
food_envir	Food environment index (0 = worst, 10 = best) (USDA Map the Meal 2014)
housing_prob	% of households with at least 1 of 4 severe problems (Comp. Housing Affordability Strategy 2009-13)
drive_alone	% of the workforce that drives alone to work (ACS 2011-15)
rural_cat	from % rural (0-20: Urban, 20.01-50: Suburban, 50.01+: Rural; Census 2015)
race_cat	from 100 - % Non-Hispanic White: (> 20: High, 10-20: Middle, 5-10: Low, < 5: Very Low; Census 2015)

Selecting Our Sample of Variables, Observations

```
ourbatch <- count_2017 %>%  
  filter(complete.cases(yrs_lost_rate),  
         state %in% c("Illinois", "Indiana", "Michigan",  
                      "Minnesota", "Ohio", "Wisconsin",  
                      "Iowa", "Missouri")) %>%  
  select(fips, state, county, yrs_lost_rate,  
         population, not_english, hh_income,  
         income_ratio, health_costs, social_assoc,  
         pct_smokers, food_envir, housing_prob,  
         drive_alone, rural_cat, race_cat)  
  
dim(ourbatch)
```

```
[1] 734 16
```

Missing Data?

```
pct_miss(ourbatch)
```

```
[1] 0
```

```
ourbatch %>% select(fips, state, county) %>%  
  summarize_all(list(n_distinct))
```

```
# A tibble: 1 x 3  
  fips state county  
  <int> <int> <int>  
1   734     8   486
```

- 1 Which states have the most/least counties?
- 2 Why do we have fewer county names than fips codes?

Counties in each state?

```
ourbatch %>% tabyl(state) %>% adorn_pct_formatting()
```

state	n	percent
Illinois	102	13.9%
Indiana	92	12.5%
Iowa	99	13.5%
Michigan	82	11.2%
Minnesota	85	11.6%
Missouri	114	15.5%
Ohio	88	12.0%
Wisconsin	72	9.8%

Across the US, the most common county name is Washington.

```
ourbatch %>% count(county, sort = TRUE) %>% filter(n > 5)
```

```
# A tibble: 8 x 2
  county      n
  <chr>    <int>
1 Jackson      8
2 Crawford     7
3 Monroe       7
4 Washington   7
5 Cass         6
6 Clinton      6
7 Jefferson    6
8 Wayne        6
```

Basic Data Summaries

Available approaches include:

- `summary`
- `mosaic` package's `inspect()`
- `skimr` package's `skim_without_charts()`
- `Hmisc` package's `describe()`

but none of them fit well on a single slide.

```
skimmed_predictors <- ourbatch %>%  
  select(-fips, -state, -county, -yrs_lost_rate) %>%  
  skim_without_charts()
```

The Quantitative Predictors

- Any concerns regarding the ranges of these variables?

```
yank(skimmed_predictors, "numeric") %>%  
  select(skim_variable, p0, p50, p100, mean, sd) %>%  
  kable(digits = 1)
```

skim_variable	p0	p50	p100	mean	sd
population	3694.0	28793.0	5238216.0	83750.2	249815.3
not_english	0.0	0.5	11.7	0.8	1.1
hh_income	28.2	49.9	97.7	50.5	9.7
income_ratio	3.2	4.1	7.2	4.2	0.5
health_costs	6387.4	9114.0	13702.9	9179.4	1087.9
social_assoc	5.2	14.5	43.7	15.5	5.6
pct_smokers	12.2	17.3	35.7	17.9	2.9
food_envir	3.8	7.7	9.3	7.6	0.8
housing_prob	6.0	12.7	24.3	12.8	3.0
drive_alone	51.7	81.5	89.2	80.9	4.5

Cuyahoga County's Values

```
cuya <- ourbatch %>% filter(county == "Cuyahoga")
```

fips	state	county	yrs_lost_rate	population	not_english
39035	Ohio	Cuyahoga	7827.8	1255921	1.85

hh_income	income_ratio	health_costs	social_assoc	pct_smokers
45.51	5.63	10023.93	9.2	18.7

food_envir	housing_prob	drive_alone	rural_cat	race_cat
6.5	18.83	80.23	Urban	High

The Categorical Predictors

```
ourbatch %>% tabyl(rural_cat, race_cat)
```

rural_cat	High	Low	Middle	Very Low
Rural	16	179	63	196
Suburban	19	87	84	10
Urban	46	2	32	0

- There are two problems I see in this table.

Oops, better re-order that race_cat factor.

```
ourbatch <- ourbatch %>%  
  mutate(race_cat = fct_relevel(race_cat,  
                                "High", "Middle", "Low"),  
         rural_cat = fct_relevel(rural_cat,  
                                "Urban", "Suburban"))  
  
ourbatch %>% tabyl(rural_cat, race_cat)
```

rural_cat	High	Middle	Low	Very Low
Urban	46	32	2	0
Suburban	19	84	87	10
Rural	16	63	179	196

- Now, would an interaction between rural_cat and race_cat be helpful in building models?

Smallest Counties by Population

```
ourbatch %>% arrange(population) %>%  
  select(state, county, population, yrs_lost_rate) %>%  
  slice(1:10) %>% kable(digits = 3)
```

state	county	population	yrs_lost_rate
Missouri	Mercer	3694	6977.9
Iowa	Adams	3796	6717.5
Missouri	Knox	3910	10133.0
Minnesota	Red Lake	4055	3980.0
Illinois	Hardin	4135	13162.8
Illinois	Pope	4226	10811.1
Minnesota	Kittson	4424	6937.3
Missouri	Schuyler	4436	6670.6
Wisconsin	Florence	4464	6281.9
Missouri	Holt	4484	7290.9


Largest Counties by Population

```
ourbatch %>% arrange(desc(population)) %>%  
  select(state, county, population, yrs_lost_rate) %>%  
  slice(1:10) %>% kable(digits = 3)
```


state	county	population	yrs_lost_rate
Illinois	Cook	5238216	6572.4
Michigan	Wayne	1759335	9960.6
Ohio	Cuyahoga	1255921	7827.8
Ohio	Franklin	1251722	7484.5
Michigan	Oakland	1242304	5805.6
Minnesota	Hennepin	1223149	4933.1
Missouri	St. Louis	1003362	6750.1
Wisconsin	Milwaukee	957735	7976.7
Indiana	Marion	939020	8992.1
Illinois	DuPage	933736	4059.0

Distribution of Population (code to follow)


Population, by County, on Linear Scale



Population, by County, plotted on Log Scale



Log (base 10) of Population Count, by County, plotted on Linear Scale



Code for plots of population distribution (start)

```
p1 <- ggplot(ourbatch, aes(x = "",  
                           y = population)) +  
  geom_violin(fill = "red") +  
  geom_boxplot(width = 0.3) + coord_flip() +  
  labs(title = "Population, by County, on Linear Scale",  
        x = "")  
  
p2 <- ggplot(ourbatch, aes(x = "",  
                           y = population)) +  
  geom_violin(fill = "pink") +  
  geom_boxplot(width = 0.3) + coord_flip() +  
  scale_y_log10() +  
  labs(title = "Population, by County, plotted on Log Scale",  
        x = "")
```

Code for plots of population distribution (finish)

```
p3 <- ggplot(ourbatch, aes(x = "",  
                           y = log10(population))) +  
  geom_violin(fill = "green") +  
  geom_boxplot(width = 0.3) + coord_flip() +  
  labs(title = "Log (base 10) of Population Count, by County",  
        x = "")  
  
p1 / p2 / p3
```


I'm going to transform the population information

```
ourbatch <- ourbatch %>%  
  mutate(log_pop = log10(population))  
  
ourbatch %$% Hmisc::describe(log_pop)
```

log_pop

n	missing	distinct	Info	Mean	Gmd
734	0	729	1	4.525	0.5612
.05	.10	.25	.50	.75	.90
3.807	3.947	4.173	4.459	4.802	5.228
.95					
5.468					

lowest : 3.567497 3.579326 3.592177 3.607991 3.616476
highest: 6.094228 6.097508 6.098962 6.245349 6.719183

Our “Model Selection” Problem

We want to build a model to effectively predict our outcome `yrs_lost_rate` using the 10 quantitative candidate predictors we've identified. (We'll add the categorical predictors in later.)

- 1 `log_pop`
- 2 `not_english`
- 3 `hh_income`
- 4 `income_ratio`
- 5 `health_costs`
- 6 `social_assoc`
- 7 `pct_smokers`
- 8 `food_envir`
- 9 `housing_prob`
- 10 `drive_alone`

We have 734 counties (observations) to use here.

Our outcome is Age-Adjusted Years Lost Rate


Code for plotting using patchwork

```
p1 <- ggplot(ourbatch, aes(yrs_lost_rate)) +  
  geom_histogram(bins = 15,  
                 fill = "navy", col = "white") +  
  labs(title = "Histogram of Years Lost, by County")  
  
p2 <- ggplot(ourbatch, aes(sample = yrs_lost_rate)) +  
  geom_qq(col = "navy") + geom_qq_line(col = "red") +  
  labs(title = "Normal Q-Q of Years Lost")  
  
p1 + p2 +  
  plot_annotation(title = "Age-Adjusted Years Lost before Ag
```


Our outcome is Age-Adjusted Years Lost Rate

Age-Adjusted Years Lost before Age 75 per 100,000 population

Histogram of Years Lost, by County



Normal Q-Q of Years Lost



yrs_lost_rate: Ten Highest Counties

state	county	population	yrs_lost_rate
Missouri	Pemiscot	17482	16011.2
Wisconsin	Menominee	4573	14634.6
Indiana	Scott	23744	13770.0
Missouri	Dunklin	30895	13442.6
Illinois	Hardin	4135	13162.8
Missouri	Iron	10125	12899.5
Minnesota	Mahnomen	5457	12186.4
Ohio	Pike	28217	12091.3
Iowa	Monona	8979	12045.0
Indiana	Fayette	23434	11965.2

yrs_lost_rate: Ten Lowest Counties

state	county	population	yrs_lost_rate
Minnesota	Houston	18773	3115.1
Minnesota	Stevens	9796	3402.8
Minnesota	Carver	98741	3408.3
Minnesota	Scott	141660	3513.7
Iowa	Lyon	11745	3648.1
Minnesota	Marshall	9423	3794.8
Indiana	Hamilton	309697	3826.8
Iowa	Sioux	34937	3848.3
Iowa	Winneshiek	20709	3867.2
Iowa	Dallas	80133	3909.7

Hmisc::describe on yrs_lost_rate

```
ourbatch %$% Hmisc::describe(yrs_lost_rate)
```

yrs_lost_rate

n	missing	distinct	Info	Mean	Gmd
734	0	732	1	7124	2004
.05	.10	.25	.50	.75	.90
4505	4969	5821	6939	8129	9415
.95					
10366					

lowest : 3115.1 3402.8 3408.3 3513.7 3648.1
highest: 13162.8 13442.6 13770.0 14634.6 16011.2

A “Kitchen Sink” Model?


Predict yrs_lost_rate using all 12 candidate predictors.

```
transformation_check <-  
  lm(yrs_lost_rate ~  
      log_pop + not_english + hh_income +  
      income_ratio + health_costs + social_assoc +  
      pct_smokers + food_envir + housing_prob +  
      drive_alone + rural_cat + race_cat,  
      data = ourbatch)  
  
glance(transformation_check) %>%  
  select(r.squared, adj.r.squared)
```

```
# A tibble: 1 x 2  
  r.squared adj.r.squared  
    <dbl>         <dbl>  
1    0.618         0.610
```


Box-Cox plot: Outcome transformation?


```
boxCox(transformation_check)
```




Logarithm of our outcome rate?

Logarithm of Age-Adjusted Years Lost before Age 75 per 100,000 population

log(Years Lost Rate) by County



Normal Q-Q of log(Years Lost)




Kitchen Sink fit to $\log(\text{yrs_lost_rate})$

```
ks_log <-  
  lm(log(yrs_lost_rate) ~  
      log_pop + not_english + hh_income +  
      income_ratio + health_costs + social_assoc +  
      pct_smokers + food_envir + housing_prob +  
      drive_alone + rural_cat + race_cat,  
      data = ourbatch)
```


Quick Look at Residual Plots from ks_log?

```
plot(ks_log, which = 1)
```



Quick Look at Residual Plots from ks_log?

```
plot(ks_log, which = 2)
```




All Four Diagnostic Plots (code)

```
par(mfrow = c(2,2))  
plot(ks_log)  
par(mfrow=c(1,1))
```


Resulting Plots on next slide...

Regression Diagnostics for ks_log


Residuals vs Fitted




Normal Q-Q



Scale-Location





Residuals vs Leverage



How about a Square Root instead?

Square Root of Age-Adjusted Years Lost before Age 75 per 100,000 population




Kitchen Sink fit to $\sqrt{\text{yrs_lost_rate}}$

```
ks_sqrt <-  
  lm(sqrt(yrs_lost_rate) ~  
    log_pop + not_english + hh_income +  
    income_ratio + health_costs + social_assoc +  
    pct_smokers + food_envir + housing_prob +  
    drive_alone + rural_cat + race_cat,  
    data = ourbatch)
```


Quick Look at Residual Plots from ks_sqrt?

```
plot(ks_sqrt, which = 1)
```







Quick Look at Residual Plots from ks_sqrt?

```
plot(ks_sqrt, which = 2)
```



All 4 Regression Diagnostic Plots for ks_sqrt



I'm going to go with the logarithmic transformation

```
ourbatch <- ourbatch %>%  
  mutate(log_yrslost = log(yrs_lost_rate))  
  
ourbatch %$% Hmisc::describe(log_yrslost)
```

log_yrslost

n	missing	distinct	Info	Mean	Gmd
734	0	732	1	8.84	0.2849
.05	.10	.25	.50	.75	.90
8.413	8.511	8.669	8.845	9.003	9.150
.95					
9.246					

lowest : 8.044017 8.132354 8.133969 8.164425 8.201962
highest: 9.485150 9.506184 9.530248 9.591144 9.681044

Our NEW “Model Selection” Problem

Let's focus for the moment on building a model to effectively predict our transformed outcome `log_yrslost` using the 10 quantitative candidate predictors we've identified. (We'll add in the categorical predictors later.)

- 1 `log_pop`
- 2 `not_english`
- 3 `hh_income`
- 4 `income_ratio`
- 5 `health_costs`
- 6 `social_assoc`
- 7 `pct_smokers`
- 8 `food_envir`
- 9 `housing_prob`
- 10 `drive_alone`

We still have 734 counties (observations).

Scatterplot Matrix?

We've got 10 quantitative predictors here, plus a (transformed) outcome.

Next, we'll show `ggpairs()` results (from `GGally` package)


```
ggpairs(ourbatch, columns = c(17, 6:9, 18))  
ggpairs(ourbatch, columns = c(10:14, 18))
```

- Run these with `message = FALSE`.
- Why include variable 18 twice, and at the end?
 - Variable order can be seen with `glimpse` or `names`.


```
names(ourbatch)
```

[1]	"fips"	"state"	"county"
[4]	"yrs_lost_rate"	"population"	"not_english"
[7]	"hh_income"	"income_ratio"	"health_costs"
[10]	"social_assoc"	"pct_smokers"	"food_envir"
[13]	"housing_prob"	"drive_alone"	"rural_cat"
[16]	"race_cat"	"log_pop"	"log_yrslost"

Scatterplot Matrix with GGally, Part I



Scatterplot Matrix with GGally, Part II



Model Selection and “Best Subsets”

Stepwise Regression: Why It's a Bad Idea

- We have lots of candidate predictors.
- We don't have an especially large sample size.

It's tempting to apply “stepwise regression” to eliminate some of the predictors and produce a more parsimonious model. This would involve:

- deciding whether we would do forwards selection, backwards elimination or a combination
- using an R function called `step` to identify changes which minimize AIC at each step

For instance, we could start with a full model using all candidate predictors, and ask R to remove predictors, one at a time, whose removal most improved (reduced) the AIC.

- When R hit the point when removing another predictor would not improve AIC, it would stop and tell us what the “winning” model was.

So, what's the problem?

Stepwise Regression is a Bad Idea

Stepwise regression **encourages you not to think.**

- Most of the time, when you collect data, you have a reason to want to use it in your model.
- This algorithm's decisions about what variables to throw away / include doesn't actually minimize AIC over the set of possible models.

In addition, there are many more substantial problems (see Harrell 2001.) Models selected using stepwise regression look better than they are. . .

- 1 The R^2 values you get by doing this are biased high.
- 2 The parameter estimates are biased away from 0, their standard errors are too small, the confidence intervals around the estimates are too narrow, and the p values are too low.
- 3 Whatever collinearity issues you have in the data are made worse.

“All Subsets” or “Best Subsets”

A somewhat more plausible strategy is to consider all possible subsets of predictors, and search in a smart way for models which are good candidates because they optimize some combination of model fit summaries.

- The `leaps` package in R has a function called `regsubsets` to help us with this approach, although we'll see that there are some problems with it, especially if we have correlated predictors.

First, we specify a “full model” to `regsubsets` and list the maximum number of predictors we are willing to include in our models, with the `nvmax` parameter.

- `regsubsets` creates a subset selection object that identifies (by exhaustive search) the best models containing 1 predictor, then 2, then 3 and so forth up to `nvmax`.

Using regsubsets in leaps

- Here, we have 734 counties, and I'll consider models with up to 8 inputs.

```
rs_models <-  
  regsubsets(log_yrslost ~ log_pop + not_english +  
             hh_income + income_ratio + health_costs +  
             social_assoc + pct_smokers + food_envir +  
             housing_prob + drive_alone,  
             data = ourbatch, nvmax = 8, nbest = 1)
```

- Other options available in regsubsets include
 - fitting more than one model per predictor count, via nbest
- The next slide displays rs_models.

Subset selection object

```
Call: regsubsets.formula(log_yrslost ~ log_pop + not_english +  
  income_ratio + health_costs + social_assoc + pct_smokers +  
  food_envir + housing_prob + drive_alone, data = ourbatch,  
  nvmax = 8, nbest = 1)
```

10 Variables (and intercept)

Forced in Forced out

log_pop	FALSE	FALSE
not_english	FALSE	FALSE
hh_income	FALSE	FALSE
income_ratio	FALSE	FALSE
health_costs	FALSE	FALSE
social_assoc	FALSE	FALSE
pct_smokers	FALSE	FALSE
food_envir	FALSE	FALSE
housing_prob	FALSE	FALSE
drive_alone	FALSE	FALSE

1 subsets of each size up to 8

Selection Algorithm: exhaustive

Summarizing Chosen Subsets

I've built a cleaner version of this on the next slide, by hand.

```
rs_summary <- summary(rs_models)
t(rs_summary$which)
```

	1	2	3	4	5	6	7
(Intercept)	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
log_pop	FALSE	FALSE	FALSE	FALSE	FALSE	TRUE	TRUE
not_english	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
hh_income	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
income_ratio	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
health_costs	FALSE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
social_assoc	FALSE	FALSE	FALSE	FALSE	TRUE	TRUE	TRUE
pct_smokers	FALSE	FALSE	TRUE	TRUE	TRUE	TRUE	TRUE
food_envir	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	TRUE
housing_prob	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
drive_alone	FALSE	FALSE	FALSE	TRUE	TRUE	TRUE	TRUE

Best Subsets from `rs_summary`

Model	Predictors
1	hh_income
2	hh_income, health_costs
3	Model 2 + pct_smokers
4	Model 3 + drive_alone
5	Model 4 + social_assoc
6	Model 5 + log_pop
7	Model 6 + food_envir
8	Model 7 + housing_prob

- `income_ratio` and `not_english` did not make any of these models
- As it turns out, in this situation, each model adds a predictor to the previous one, but that doesn't have to be the case.

Obtaining Fit Quality Statistics

```
rs_winners <- tbl_df(rs_summary$which) %>%  
  mutate(inputs = 1:(rs_models$nvmax - 1),  
         r2 = rs_summary$rsq,  
         adjr2 = rs_summary$adjr2,  
         cp = rs_summary$cp,  
         bic = rs_summary$bic,  
         rss = rs_summary$rss) %>%  
  select(inputs, adjr2, cp, bic, everything())
```

Looking at rs_winners

```
rs_winners %>% slice(1:4)
```

```
# A tibble: 4 x 17
```

```
  inputs adjr2      cp    bic `(Intercept)` log_pop
    <int> <dbl>  <dbl> <dbl> <lgl>          <lgl>
1       1 0.467 297.   -450. TRUE          FALSE
2       2 0.564 112.   -591. TRUE          FALSE
3       3 0.597  48.1  -644. TRUE          FALSE
4       4 0.620   5.06 -681. TRUE          FALSE
# ... with 11 more variables: not_english <lgl>,
#   hh_income <lgl>, income_ratio <lgl>,
#   health_costs <lgl>, social_assoc <lgl>,
#   pct_smokers <lgl>, food_envir <lgl>,
#   housing_prob <lgl>, drive_alone <lgl>, r2 <dbl>,
#   rss <dbl>
```

Suppose we plot summary statistics for each model.

Start with R^2 and residual sum of squares for the models.

```
rs_winners %>% select(inputs, r2, rss) %>% kable(digits = 4)
```

inputs	r2	rss
1	0.4678	24.9322
2	0.5647	20.3913
3	0.5989	18.7913
4	0.6222	17.6989
5	0.6236	17.6316
6	0.6242	17.6031
7	0.6249	17.5716
8	0.6251	17.5652


- Which model will look best, by R^2 ? By rss?
- How should we plot these results?

Plotting Raw R^2 values (code)

```
ggplot(rs_winners, aes(x = inputs, y = r2,  
                        label = round(r2,3))) +  
  geom_line() +  
  geom_label() +  
  geom_label(data = subset(rs_winners,  
                           r2 == max(r2)),  
             aes(x = inputs, y = r2,  
                 label = round(r2,3)),  
             fill = "black", col = "white") +  
  labs(x = "# of regression inputs",  
       y = "R-squared")
```

What will this do?

Plotting Raw R^2 and Residual SS values



- These plots will always select the largest model available. Why?

So, which Subsets Look Best?


The `regsubsets` summary includes three model quality statistics:

- Adjusted R^2 , which we'd like to maximize
- Mallows' C_p , which we'd like to minimize
- BIC, the Bayes Information Criterion, which we'd like to minimize

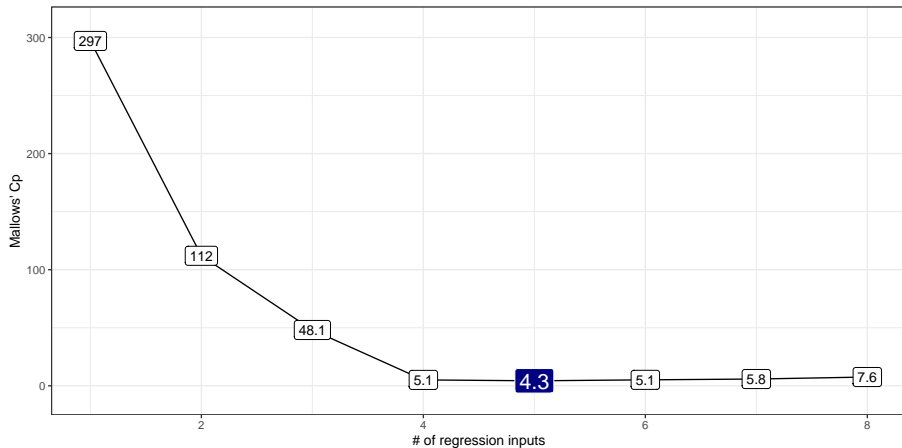
Plotting Adjusted R^2 values

```
ggplot(rs_winners, aes(x = inputs, y = adjr2,  
                        label = round(adjr2,3))) +  
  geom_line() +  
  geom_label() +  
  geom_label(data = subset(rs_winners,  
                           adjr2 == max(adjr2)),  
            aes(x = inputs, y = adjr2,  
                label = round(adjr2,3)),  
            fill = "yellow", col = "blue", size = 6) +  
  scale_y_continuous(expand = expand_scale(mult = .1)) +  
  labs(x = "# of regression inputs",  
       y = "Adjusted R-squared")
```



Adjusted R^2 values for our subsets



Mallows' C_p for our subsets



BIC for our subsets



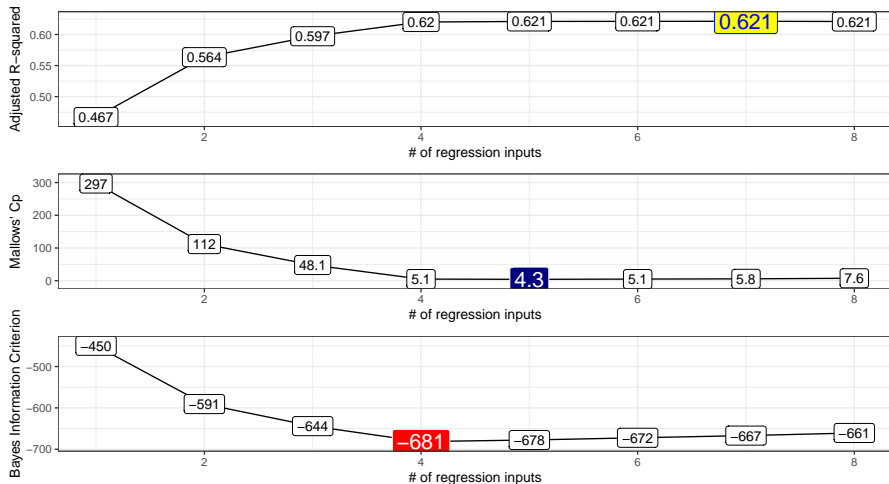
So, which Subsets Look Best?

We can use `which.max` and `which.min` to identify from the subset list the models selected by each of the three summaries we've been plotting...

```
tibble(AdjR2 = which.max(rs_summary$adjr2),  
       Cp = which.min(rs_summary$cp),  
       BIC = which.min(rs_summary$bic))
```

```
# A tibble: 1 x 3  
  AdjR2    Cp    BIC  
  <int> <int> <int>  
1     7     5     4
```

By the plots, we have ...



Coefficients for Candidate Models

```
coef(rs_models, id = 4)
```

(Intercept)	hh_income	health_costs	pct_smokers
7.899140e+00	-1.310119e-02	4.882626e-05	2.233060e-02
drive_alone			
9.329765e-03			

```
coef(rs_models, id = 5)
```

(Intercept)	hh_income	health_costs	social_assoc
8.034529e+00	-1.370597e-02	4.904949e-05	-1.990994e-03
pct_smokers	drive_alone		
1.991591e-02	8.924561e-03		

Coefficients for Candidate Models (continued)

```
coef(rs_models, id = 7)
```

(Intercept)	log_pop	hh_income	health_costs
8.189596e+00	-1.962619e-02	-1.293658e-02	4.946928e-05
social_assoc	pct_smokers	food_envir	drive_alone
-2.693840e-03	1.873528e-02	-1.188579e-02	9.089601e-03

Our Candidate Models

Inputs	Predictors
--------	------------

- | | |
|---|---|
| 4 | hh_income, health_costs, pct_smokers, drive_alone |
| 5 | Model 4 + social_assoc |
| 6 | Model 5 + log_pop and food_envir |
-

In-Sample Comparisons of our Candidate Models

```
m4 <- ourbatch %$%  
  lm(log_yrslost ~ hh_income + health_costs +  
      pct_smokers + drive_alone)  
  
m5 <- ourbatch %$%  
  lm(log_yrslost ~ hh_income + health_costs +  
      pct_smokers + drive_alone + social_assoc)  
  
m7 <- ourbatch %$%  
  lm(log_yrslost ~ hh_income + health_costs +  
      pct_smokers + drive_alone + social_assoc +  
      log_pop + food_envir)
```

Models are **nested** so comparisons within samples are straightforward.

Comparing the in-sample fits of the models

```
bind_rows(glance(m4), glance(m5), glance(m7)) %>%  
  mutate(model = c("m4", "m5", "m7")) %>%  
  select(model, r2 = r.squared, adjr2 = adj.r.squared,  
         AIC, BIC, sigma) %>%  
  kable(digits = c(0, 3, 4, 1, 1, 4))
```

model	r2	adjr2	AIC	BIC	sigma
m4	0.622	0.6201	-639.2	-611.6	0.1558
m5	0.624	0.6211	-639.9	-607.8	0.1556
m7	0.625	0.6213	-638.5	-597.1	0.1556

Comparisons in-sample with anova

```
anova(m4, m5, m7)
```

Analysis of Variance Table

Model 1: log_yrslost ~ hh_income + health_costs + pct_smokers

Model 2: log_yrslost ~ hh_income + health_costs + pct_smokers
social_assoc

Model 3: log_yrslost ~ hh_income + health_costs + pct_smokers
social_assoc + log_pop + food_envir

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	729	17.699				
2	728	17.632	1	0.067272	2.7794	0.09591 .
3	726	17.572	2	0.059984	1.2392	0.29024

Signif. codes:

0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

What if the models you're comparing aren't nested?

What if you're comparing:

- Model A: `lm(y = x1 + x2 + x3, data = dataset)`
- Model B: `lm(y = x1 + x4 + x5, data = dataset)`

Then ...

- default p values from the ANOVA table comparing Model A to Model B aren't reasonable
- AIC and BIC are OK, can also use adjusted R^2 to help make a decision within the model building sample
- Still useful to think about out-of-sample prediction and cross-validation

Cross-Validation by Validation Split

A First Cross-Validation Approach

We'll start with a validation split. . .

- 1 Partition the data into a training set and a test set.
 - We'll identify a random sample of 80% of the ourbatch data into our training sample, and the rest will be in our test sample.
- 2 Build (train) the three candidate models we identified with “best subsets” by running them on the training set.
 - Note that I actually used the whole sample (not just the training sample) to perform best subsets here. I'm not claiming that was a good idea, just demonstrating.
- 3 Evaluate the model by looking at its effectiveness in the test set.
 - In this case, I'll compare the validation results for the three models we identified.

Measures of Fit Quality used in Cross-Validation

- 1 R-squared (R^2): the squared correlation between the observed outcome values and the predicted values by the model.
- 2 Root Mean Squared Error (RMSE), which measures the average prediction error made by the model in predicting the outcome for an observation. That is, the average difference between the observed known outcome values and the values predicted by the model. The lower the RMSE, the better the model.
- 3 Mean Absolute Error (MAE), an alternative to the RMSE that is less sensitive to outliers. It corresponds to the average absolute difference between observed and predicted outcomes. The lower the MAE, the better the model.

See the link on the Class 7 README for more details.

Validation Split Approach

We'll use the `createDataPartition` function from the `caret` package.

```
set.seed(432)
training.samples <- ourbatch$log_yrslost %>%
  createDataPartition(p = 0.8, list = FALSE)

our_train <- ourbatch[training.samples,]
our_test <- ourbatch[-training.samples,]
```

Remember that our models are built to predict `log_yrs_lost = log(yrs_lost_rate)`, but we actually want to predict the `yrs_lost_rate`, so we'll need to do some back-transformation after making predictions.

Validation Split Approach for m4

- 1 Run the m4 model in the training sample.

```
m4_train <- lm(log_yrslost ~ hh_income + health_costs +  
               pct_smokers + drive_alone,  
               data = our_train)
```

- 2 Obtain predicted ($\log(\text{yrs_lost_rate})$) values in the test sample.

```
m4_pred_logs <- m4_train %>% predict(our_test)
```

- 3 Exponentiate the predicted ($\log(\text{yrs_lost_rate})$) values to get predicted yrs_lost_rate values within the test sample.

```
m4_pred_yrslost <- exp(m4_pred_logs)
```

Validation Split Approach for m4 (continued)

- ④ Use the R2, RMSE and MAE functions from the caret package to calculate key summaries of the predictions we made for yrs_lost_rate using our model 4, and how they compare to the observed yrs_lost_rate values in the training sample.

```
m4_summaries <- tibble(  
  model = "m4",  
  R2 = R2(m4_pred_yrslost, our_test$yrs_lost_rate),  
  RMSE = RMSE(m4_pred_yrslost, our_test$yrs_lost_rate),  
  MAE = MAE(m4_pred_yrslost, our_test$yrs_lost_rate),  
  pred_err_rate = RMSE / mean(our_test$yrs_lost_rate)  
)
```

Validation Split for Model m4: Results

```
m4_summaries
```

```
# A tibble: 1 x 5
```

	model	R2	RMSE	MAE	pred_err_rate
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	m4	0.456	1253.	997.	0.178

Validation Split for Model m5

```
m5_train <- lm(log_yrslost ~ hh_income + health_costs +  
               pct_smokers + drive_alone + social_assoc,  
               data = our_train)  
  
m5_pred_logs <- m5_train %>% predict(our_test)  
m5_pred_yrslost <- exp(m5_pred_logs)  
  
m5_summaries <- tibble(  
  model = "m5",  
  R2 = R2(m5_pred_yrslost, our_test$yrs_lost_rate),  
  RMSE = RMSE(m5_pred_yrslost, our_test$yrs_lost_rate),  
  MAE = MAE(m5_pred_yrslost, our_test$yrs_lost_rate),  
  pred_err_rate = RMSE / mean(our_test$yrs_lost_rate)  
)
```

Validation Split for Model m7

```
m7_train <- lm(log_yrslost ~ hh_income + health_costs +  
               pct_smokers + drive_alone +  
               social_assoc + log_pop + food_envir,  
               data = our_train)  
  
m7_pred_logs <- m7_train %>% predict(our_test)  
m7_pred_yrslost <- exp(m7_pred_logs)  
  
m7_summaries <- tibble(  
  model = "m7",  
  R2 = R2(m7_pred_yrslost, our_test$yrs_lost_rate),  
  RMSE = RMSE(m7_pred_yrslost, our_test$yrs_lost_rate),  
  MAE = MAE(m7_pred_yrslost, our_test$yrs_lost_rate),  
  pred_err_rate = RMSE / mean(our_test$yrs_lost_rate)  
)
```

Validation Split Approach Model Comparison

```
bind_rows(m4_summaries, m5_summaries, m7_summaries) %>%  
  kable(digits = c(0, 3, 1, 1, 4))
```

model	R2	RMSE	MAE	pred_err_rate
m4	0.456	1252.9	996.6	0.1780
m5	0.457	1250.8	998.5	0.1777
m7	0.458	1250.3	996.9	0.1776

- Which of these models shows the strongest performance?

Validation Split Approach Model Comparison

```
bind_rows(m4_summaries, m5_summaries, m7_summaries) %>%  
  kable(digits = c(0, 3, 1, 1, 4))
```

model	R2	RMSE	MAE	pred_err_rate
m4	0.456	1252.9	996.6	0.1780
m5	0.457	1250.8	998.5	0.1777
m7	0.458	1250.3	996.9	0.1776

- In validation R^2 , we're looking at the squared correlation coefficient for the predictions and the observed values of `yrs_lost_rate`, so we want larger numbers.
 - Note that these validated R^2 values are not comparable to our earlier R^2 values, because now we're trying to predict `yrs_lost_rate` rather than `log(yrs_lost_rate)`.
 - Here, m7 is largest, followed by m5 and m4.

Validation Split Approach Model Comparison

```
bind_rows(m4_summaries, m5_summaries, m7_summaries) %>%  
  kable(digits = c(0, 3, 1, 1, 4))
```

model	R2	RMSE	MAE	pred_err_rate
m4	0.456	1252.9	996.6	0.1780
m5	0.457	1250.8	998.5	0.1777
m7	0.458	1250.3	996.9	0.1776

- The Root Mean Square Errors (RMSE) are similar, with m7 slightly smaller (better) than m5 and m4.
- The Mean Absolute Errors (MAE) in the validation sample are also similar, with m4 now a little smaller (better) than m7 or m5.
- The prediction error rate is just the RMSE divided by the mean of the actual yrs_lost_rate. No new information here on model fit.

K-Fold Cross Validation

Algorithm for K-fold Cross-Validation

- ➊ Randomly split the data set into k subsets (for k -fold validation)
- ➋ Reserve one subset and train the model on all other subsets
- ➌ Test the model on the reserved subset and record the prediction error
- ➍ Repeat this process until each of the k subsets has served as the test set.
- ➎ Compute the average of the k recorded errors. This is called the cross-validation error serving as the performance metric for the model.

Usually k is 5 or 10.

10-fold cross validation for model m4

```
# Define training control
set.seed(43201)
train.control <- trainControl(method = "cv", number = 10)
# Train the model
model4_cv <- train(log_yrslost ~ hh_income + health_costs +
                    pct_smokers + drive_alone,
                    data = ourbatch, method = "lm",
                    trControl = train.control)
```

Summarize the results from model4_cv

```
model4_cv
```

Linear Regression

734 samples

4 predictor

No pre-processing

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 661, 659, 660, 659, 661, 661, ...

Resampling results:

RMSE	Rsquared	MAE
0.1555325	0.6284756	0.1218445

Tuning parameter 'intercept' was held constant at a value of TRUE

Model Fit by model4_cv summary

```
tidy(summary(model4_cv)) %>% kable(digits = 3)
```

term	estimate	std.error	statistic	p.value
(Intercept)	7.899	0.124	63.698	0
hh_income	-0.013	0.001	-16.199	0
health_costs	0.000	0.000	8.256	0
pct_smokers	0.022	0.003	8.093	0
drive_alone	0.009	0.001	6.708	0

Model Fit by model4_cv summary

```
glance(summary(model4_cv)) %>% kable(digits = 3)
```

r.squared	adj.r.squared	sigma	statistic	p.value	df
0.622	0.62	0.156	300.152	0	5

But, of course, these summaries are on the log scale again.

10-fold cross validation for model m5

```
# Define training control
set.seed(43202)
train.control <- trainControl(method = "cv", number = 10)
# Train the model
model5_cv <- train(log_yrslost ~ hh_income + health_costs +
                    pct_smokers + drive_alone +
                    social_assoc,
                    data = ourbatch, method = "lm",
                    trControl = train.control)
```

Summarize the results from model5_cv

```
model5_cv
```

Linear Regression

734 samples

5 predictor

No pre-processing

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 660, 661, 660, 661, 660, 661, ...

Resampling results:

RMSE	Rsquared	MAE
0.1557501	0.6272377	0.1222321

Tuning parameter 'intercept' was held constant at a value of TRUE

10-fold cross validation for model m7

```
# Define training control
set.seed(43203)
train.control <- trainControl(method = "cv", number = 10)
# Train the model
model7_cv <- train(log_yrslost ~ hh_income + health_costs +
                    pct_smokers + drive_alone +
                    social_assoc + log_pop + food_envir,
                    data = ourbatch, method = "lm",
                    trControl = train.control)
```

Summarize the results from model7_cv

```
model7_cv
```

Linear Regression

734 samples

7 predictor

No pre-processing

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 660, 658, 662, 659, 662, 660, ...

Resampling results:

RMSE	Rsquared	MAE
0.1567999	0.6207551	0.122406

Tuning parameter 'intercept' was held constant at a value of TRUE

Comparing 10-fold Cross-Validation Results

Cross-validation results, predicting log_yrs_lost.

```
res_cv10 <- bind_rows(  
  model4_cv$results,  
  model5_cv$results,  
  model7_cv$results) %>%  
  mutate(model = c("m4", "m5", "m7")) %>%  
  select(model, RMSE, Rsquared, MAE)  
  
res_cv10 %>% kable(digits = c(0, 4, 3, 4))
```

model	RMSE	Rsquared	MAE
m4	0.1555	0.628	0.1218
m5	0.1558	0.627	0.1222
m7	0.1568	0.621	0.1224

- Which model looks best?





Fitted model4 on full data

Let's look at diagnostic plots, on the full sample...

```
m4_full <- lm(log_yrslost ~ hh_income + health_costs +  
              pct_smokers + drive_alone,  
              data = ourbatch)  
  
par(mfrow=c(2,2))  
plot(m4_full)  
par(mfrow = c(1,1))
```

See next slide for results...

Diagnostics: model14 on full data



Back-Transformed In-Sample Predictions

If we want to make predictions on the original `yrs_lost_rate` scale, then we could back-transform the predicted values manually, just to describe results in the full sample. So this isn't a validated summary - just an in-sample comparison.

```
m4_cv_pred_yrs_lost <- exp(predict(model4_cv))
m4_cv_bt_summs <- tibble(
  model = "m4_cv",
  R2 = R2(m4_cv_pred_yrs_lost, ourbatch$yrs_lost_rate),
  RMSE = RMSE(m4_cv_pred_yrs_lost, ourbatch$yrs_lost_rate),
  MAE = MAE(m4_cv_pred_yrs_lost, ourbatch$yrs_lost_rate))
m4_cv_bt_summs
```



```
# A tibble: 1 x 4
  model      R2  RMSE  MAE
  <chr> <dbl> <dbl> <dbl>
1 m4_cv 0.622 1116.  855.
```

Back-Transformed In-Sample Predictions for `model5_cv`

```
m5_cv_pred_yrs_lost <- exp(predict(model5_cv))

m5_cv_bt_summs <- tibble(
  model = "m5_cv",
  R2 = R2(m5_cv_pred_yrs_lost, ourbatch$yrs_lost_rate),
  RMSE = RMSE(m5_cv_pred_yrs_lost, ourbatch$yrs_lost_rate),
  MAE = MAE(m5_cv_pred_yrs_lost, ourbatch$yrs_lost_rate)
)

m5_cv_bt_summs
```

A tibble: 1 x 4

	model	R2	RMSE	MAE
	<chr>	<dbl>	<dbl>	<dbl>
1	m5_cv	0.623	1115.	855.

Back-Transformed In-Sample Predictions for model7_cv

```
m7_cv_pred_yrs_lost <- exp(predict(model7_cv))

m7_cv_bt_summs <- tibble(
  model = "m7_cv",
  R2 = R2(m7_cv_pred_yrs_lost, ourbatch$yrs_lost_rate),
  RMSE = RMSE(m7_cv_pred_yrs_lost, ourbatch$yrs_lost_rate),
  MAE = MAE(m7_cv_pred_yrs_lost, ourbatch$yrs_lost_rate)
)

m7_cv_bt_summs

# A tibble: 1 x 4
  model    R2  RMSE  MAE
  <chr> <dbl> <dbl> <dbl>
1 m7_cv 0.625 1112.  851.
```


Prediction Quality of Cross-Validated Models on the Full Sample (back-transformed)

Which of the three models displays has the strongest in-sample predictions?

```
res_cv10_bt <- bind_rows(  
  m4_cv_bt_summs,  
  m5_cv_bt_summs,  
  m7_cv_bt_summs) %>%  
  select(model, R2, RMSE, MAE)  
  
res_cv10_bt %>% kable(digits = c(0, 3, 1, 1))
```

model	R2	RMSE	MAE
m4_cv	0.622	1115.9	855.2
m5_cv	0.623	1115.5	854.6
m7_cv	0.625	1111.8	851.4

Variable Importance

For a linear regression, the absolute value of the t statistic for each model parameter is used to specify the importance of each variable, scaled to have a maximum value of 100.

```
tidy(summary(model4_cv)) %>% kable(digits = c(0,6,6,2,2))
```

term	estimate	std.error	statistic	p.value
(Intercept)	7.899140	0.124008	63.70	0
hh_income	-0.013101	0.000809	-16.20	0
health_costs	0.000049	0.000006	8.26	0
pct_smokers	0.022331	0.002759	8.09	0
drive_alone	0.009330	0.001391	6.71	0

Variable Importance for model4_cv


```
varImp(model4_cv)
```

```
lm variable importance
```

	Overall
hh_income	100.00
health_costs	16.31
pct_smokers	14.59
drive_alone	0.00


Plotting Variable Importance for model4_cv

```
plot(varImp(model4_cv))
```




Plotting Variable Importance for model15_cv

```
plot(varImp(model15_cv))
```



Plotting Variable Importance for model17_cv

```
plot(varImp(model17_cv))
```



Incorporating Categorical Variables in Best Subsets

Let's start again with regsubsets.

- 1 First, we'll split into a training sample and a test sample.
- 2 Then, we'll look at models that can include the categorical predictors as well as the quantitative ones we looked at previously.
- 3 We'll use best subsets to identify 2 models for each predictor count.
- 4 But we'll restrict the number of predictors more.

Partitioning into Training and Test Samples

```
set.seed(4322020)
training_samples2 <- ourbatch$log_yrslost %>%
  createDataPartition(p = 0.7, list = FALSE)
```

```
our_train2 <- ourbatch[training.samples,]
our_test2 <- ourbatch[-training.samples,]
```

70% in training, 30% in test here.

```
dim(our_train2)
```

```
[1] 590 18
```

Setting up regsubsets in our_train2

```
rs2_models <-  
  regsubsets(log_yrslost ~ health_costs + pct_smokers +  
             food_envir + race_cat + rural_cat +  
             log_pop + income_ratio,  
             data = our_train2, nvmax = 7)
```

- Best models for 1:7 predictors (what do we expect to see?)
- Including the two factors (race_cat and rural_cat)

Resulting Subsets (any surprises?)

```
rs2_summary <- summary(rs2_models)
t(rs2_summary$which)
```

	1	2	3	4	5	6	7
(Intercept)	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
health_costs	FALSE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
pct_smokers	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
food_envir	FALSE	FALSE	TRUE	TRUE	TRUE	TRUE	TRUE
race_catMiddle	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	TRUE
race_catLow	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
race_catVery Low	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
rural_catSuburban	FALSE	FALSE	FALSE	FALSE	TRUE	TRUE	TRUE
rural_catRural	FALSE	FALSE	FALSE	FALSE	TRUE	TRUE	TRUE
log_pop	FALSE	FALSE	FALSE	TRUE	FALSE	FALSE	FALSE
income_ratio	FALSE	FALSE	FALSE	FALSE	FALSE	TRUE	TRUE

Which Subsets Look Best?



```
tibble(AdjR2 = which.max(rs2_summary$adjr2),  
       Cp = which.min(rs2_summary$cp),  
       BIC = which.min(rs2_summary$bic))
```

```
# A tibble: 1 x 3  
  AdjR2    Cp    BIC  
  <int> <int> <int>  
1     7     6     5
```

Obtaining Fit Quality Statistics

```
rs2_winners <- tbl_df(rs2_summary$which) %>%  
  mutate(inputs = 1:(rs2_models$nvmax - 1),  
         r2 = rs2_summary$rsq,  
         adjr2 = rs2_summary$adjr2,  
         cp = rs2_summary$cp,  
         bic = rs2_summary$bic,  
         rss = rs2_summary$rss) %>%  
  select(inputs, adjr2, cp, bic, everything())
```

Plots from rs2_winners



Candidate Models are 5, 6 and 7 (from plots)

```
t(rs2_summary$which)
```

	1	2	3	4	5	6	7
(Intercept)	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
health_costs	FALSE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
pct_smokers	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE	TRUE
food_envir	FALSE	FALSE	TRUE	TRUE	TRUE	TRUE	TRUE
race_catMiddle	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	TRUE
race_catLow	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
race_catVery Low	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE	FALSE
rural_catSuburban	FALSE	FALSE	FALSE	FALSE	TRUE	TRUE	TRUE
rural_catRural	FALSE	FALSE	FALSE	FALSE	TRUE	TRUE	TRUE
log_pop	FALSE	FALSE	FALSE	TRUE	FALSE	FALSE	FALSE
income_ratio	FALSE	FALSE	FALSE	FALSE	FALSE	TRUE	TRUE

- Let's get the coefficients

rs2: models 5 and 6

```
coef(rs2_models, id = 5)
```

(Intercept)	health_costs	pct_smokers
8.0899637053	0.0000638064	0.0356113385
food_envir	rural_catSuburban	rural_catRural
-0.0767955853	0.0888555989	0.1364106956

```
coef(rs2_models, id = 6)
```

(Intercept)	health_costs	pct_smokers
7.888543e+00	6.357061e-05	3.517056e-02
food_envir	rural_catSuburban	rural_catRural
-6.747313e-02	9.588149e-02	1.486048e-01
income_ratio		
3.123889e-02		

rs2: model 7

```
coef(rs2_models, id = 7)
```

(Intercept)	health_costs	pct_smokers
7.902062e+00	6.311798e-05	3.528932e-02
food_envir	race_catMiddle	rural_catSuburban
-6.758231e-02	-1.731371e-02	9.588537e-02
rural_catRural	income_ratio	
1.432441e-01	3.047778e-02	

So we have three candidate models. Cross-validate?

- 1 Run rs2 model 5 in the training sample (`our_train2`).
- 2 Obtain predicted ($\log(\text{yrs_lost_rate})$) in the test sample (`our_test2`).
- 3 Exponentiate to get predicted `yrs_lost_rate` in test sample.
- 4 Then repeat steps 1-3 for models 6 and 7.

Cross-Validating rs2 model 5

```
rs2_m5_train <-  
  lm(log_yrslost ~ pct_smokers + health_costs +  
      food_envir + rural_cat,  
      data = our_train2)  
  
rs2_m5_pred_logs <- rs2_m5_train %>% predict(our_test2)  
rs2_m5_pred_yrslost <- exp(rs2_m5_pred_logs)
```

Can use `postResample` to obtain key summaries

```
postResample(rs2_m5_pred_yrslost, our_test2$yrs_lost_rate)
```

RMSE	Rsquared	MAE
1422.1611676	0.3282872	1117.2404286

Cross-Validating rs2 model 6

```
rs2_m6_train <-  
  lm(log_yrslost ~ pct_smokers + health_costs +  
      food_envir + rural_cat + income_ratio,  
      data = our_train2)  
  
rs2_m6_pred_logs <- rs2_m6_train %>% predict(our_test2)  
rs2_m6_pred_yrslost <- exp(rs2_m6_pred_logs)  
  
postResample(rs2_m6_pred_yrslost, our_test2$yrs_lost_rate)
```

RMSE	Rsquared	MAE
1419.7607492	0.3292653	1114.3468183

Cross-Validating rs2 model 7

Notice the need for trickiness here...

```
rs2_m7_train <-  
  lm(log_yrslost ~ pct_smokers + health_costs +  
      food_envir + rural_cat + income_ratio +  
      (race_cat == "Middle"),  
      data = our_train2)  
  
rs2_m7_pred_logs <- rs2_m7_train %>% predict(our_test2)  
rs2_m7_pred_yrslost <- exp(rs2_m7_pred_logs)  
  
postResample(rs2_m7_pred_yrslost, our_test2$yrs_lost_rate)
```

RMSE	Rsquared	MAE
1420.7300343	0.3289402	1111.8836867

Did I fit the right rs2_m7_train?

```
coef(rs2_m7_train)
```

(Intercept)	pct_smokers
7.902062e+00	3.528932e-02
health_costs	food_envir
6.311798e-05	-6.758231e-02
rural_catSuburban	rural_catRural
9.588537e-02	1.432441e-01
income_ratio	race_cat == "Middle"TRUE
3.047778e-02	-1.731371e-02

```
coef(rs2_models, id = 7)
```

(Intercept)	health_costs	pct_smokers
7.902062e+00	6.311798e-05	3.528932e-02
food_envir	race_catMiddle	rural_catSuburban
-6.758231e-02	-1.731371e-02	9.588537e-02
rural_catRural	income_ratio	
1.432441e-01	3.047778e-02	

Key Summaries for rs2 model 5

```
rs2_m5_summaries <- tibble(  
  model = "rs2_m5",  
  R2 = R2(rs2_m5_pred_yrslost, our_test2$yrs_lost_rate),  
  RMSE = RMSE(rs2_m5_pred_yrslost, our_test2$yrs_lost_rate),  
  MAE = MAE(rs2_m5_pred_yrslost, our_test2$yrs_lost_rate),  
)
```

Then calculate 6 and 7 summaries in the same way (hidden here)

Cross-Validation Model Comparison

```
bind_rows(rs2_m5_summaries, rs2_m6_summaries,  
          rs2_m7_summaries) %>%  
  kable(digits = c(0, 4, 1, 1))
```

model	R2	RMSE	MAE
rs2_m5	0.3283	1422.2	1117.2
rs2_m6	0.3293	1419.8	1114.3
rs2_m7	0.3289	1420.7	1111.9

- Which model shows the strongest performance?

Obtaining Regression Diagnostics in the Full Sample


We'll look at Model 6...

```
rs2_m6 <- lm(log_yrslost ~ pct_smokers + health_costs +  
             food_envir + rural_cat + income_ratio,  
             data = ourbatch)  
  
par(mfrow = c(2,2))  
plot(rs2_m6)  
par(mfrow=c(1,1))
```


Resulting Plots on next slide...

Regression Diagnostics for rs2_m6 (Full Sample)


Residuals vs Fitted




Normal Q-Q



Scale-Location



Residuals vs Leverage



Still to Come

Other Model Selection Strategies

- Ridge Regression
- The Lasso
- Elastic Net

Considering Non-Linearity in the Predictor Set

- The Spearman ρ^2 plot as a way to identify potential variables to consider non-linear terms.
- Using Polynomials, Interaction Terms and Restricted Cubic Splines
- Using `ols` from `rms` to fit linear models

Predicting a Binary Outcome

- `bestglm` for best subsets?
- Classification Strategies

Also still to come

- Using `tidymodels` approaches
- Regression models for other types of outcomes
- Using sampling weights
- Predictor pre-processing
- Multiple Imputation as pre-processing