

432 Class 9 Slides

github.com/THOMASELOVE/2020-432

2020-02-13

Setup

```
library(janitor); library(magrittr); library(here)
```

```
library(skimr)
```

```
library(broom)
```

```
library(rms)
```

```
library(patchwork)
```

```
library(GGally)
```

```
library(tidyverse)
```

```
theme_set(theme_bw())
```

Today's Materials

- The maleptsd data
- Using `ols` to fit a linear model
 - ANOVA in `ols`
 - Plot Effects with `summary` and `Predict`
 - Validating summary statistics like R^2
 - Nomogram
 - Evaluating Calibration
- Spending Degrees of Freedom on Non-Linearity
 - The Spearman ρ^2 (rho-squared) plot
- Building Non-Linear Predictors with
 - Polynomial Functions
 - Splines, including Restricted Cubic Splines

The maleptsd data: Background and Exploration

The maleptsd data

The maleptsd file on our web site contains information on PTSD (post traumatic stress disorder) symptoms following childbirth for 64 fathers¹. There are ten predictors and the response is a measure of PTSD symptoms. The raw, untransformed values (ptsd_raw) are right skewed and contain zeros, so we will work with a transformation, specifically, $\text{ptsd} = \log(\text{ptsd_raw} + 1)$ as our outcome, which also contains a lot of zeros.

```
maleptsd <- read_csv(here("data/maleptsd.csv")) %>%  
  clean_names() %>%  
  mutate(ptsd = log(ptsd_raw + 1))
```

¹Source: Ayers et al. 2007 *J Reproductive and Infant Psychology*. The data are described in more detail in Wright DB and London K (2009) *Modern Regression Techniques Using R* Sage Publications.



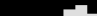








Skimming the maleptsd data

```
> maleptsd %>% select(-id, -ptsd_raw) %>% skim()
```

```
Skim summary statistics
```


```
n obs: 64
```

```
n variables: 11
```


-- Variable type:numeric -----											
variable	missing	complete	n	mean	sd	p0	p25	p50	p75	p100	hist
aff	0	64	64	8.84	3.08	0	7	9.5	11	17	
bond	0	64	64	22.52	3.07	9	21	23	24.25	28	
cons	0	64	64	48.98	11.15	0	45.75	51	55	65	
contr	0	64	64	44.2	14.84	4.5	35.1	41.75	53.98	78.5	
neg	0	64	64	21.09	11.6	0.7	11.17	20.65	30.42	45.4	
over2	0	64	64	2.8	3.34	0	0	1	5	10	
over3	0	64	64	2.72	3.13	0	0	1	5	10	
over5	0	64	64	9.12	1.34	4	9	9.5	10	10	
posit	0	64	64	35.44	11.02	2.5	27.05	37	43.35	50.1	
ptsd	0	64	64	1.59	1.27	0	0	1.61	2.74	3.95	
sup	0	64	64	13	5.87	1.2	9.28	14.25	18.3	20	

Transformation of Outcome

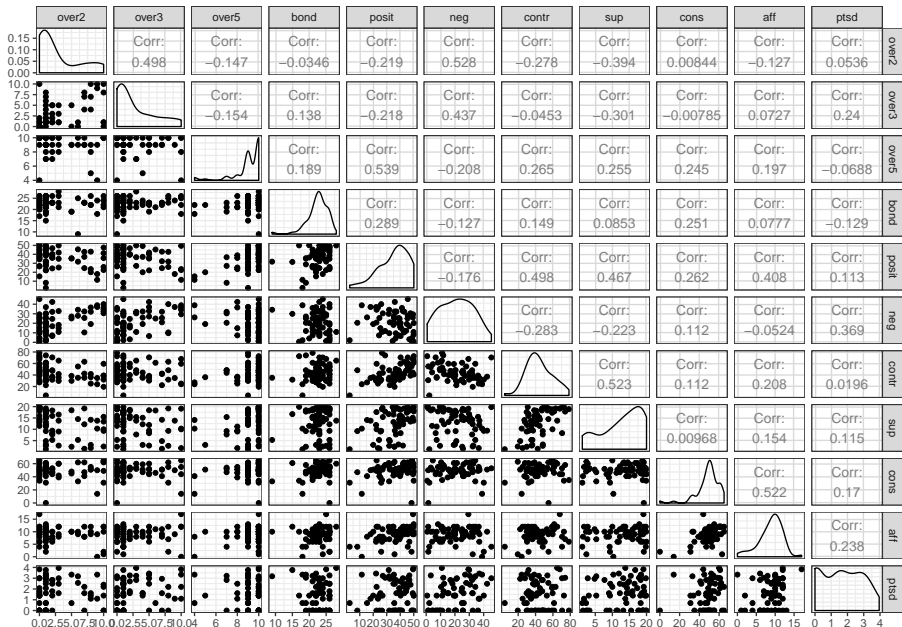
Original PTSD scores



After Transformation



Scatterplot Matrix



Using `ols` to fit a Two-Predictor Model

- `ols` is part of the `rms` package, by Frank Harrell and colleagues.

```
dd <- datadist(maleptsd)
options(datadist = "dd")

mod_first <- ols(ptsd ~ over2 + over3, data = maleptsd,
                 x = TRUE, y = TRUE)
```

Contents of mod_first?

```
> mod_first
```

Linear Regression Model

```
ols(formula = ptsd ~ over2 + over3, data = maleptsd)
```

		Model Likelihood		Discrimination	
		Ratio Test		Indexes	
Obs	64	LR chi2	4.18	R2	0.063
sigma	1.2500	d.f.	2	R2 adj	0.033
d.f.	61	Pr(> chi2)	0.1235	g	0.345

Residuals

	Min	1Q	Median	3Q	Max
	-2.259244	-1.337198	0.008866	1.140664	2.333183

	Coef	S.E.	t	Pr(> t)
Intercept	1.3733	0.2202	6.24	<0.0001
over2	-0.0333	0.0544	-0.61	0.5425
over3	0.1149	0.0579	1.98	0.0518

ANOVA for mod_first fit by ols

```
anova(mod_first)
```

Analysis of Variance

Response: ptsd

Factor	d.f.	Partial SS	MS	F	P
over2	1	0.5862441	0.5862441	0.38	0.5425
over3	1	6.1458656	6.1458656	3.93	0.0518
REGRESSION	2	6.4382526	3.2191263	2.06	0.1362
ERROR	61	95.3127887	1.5625047		

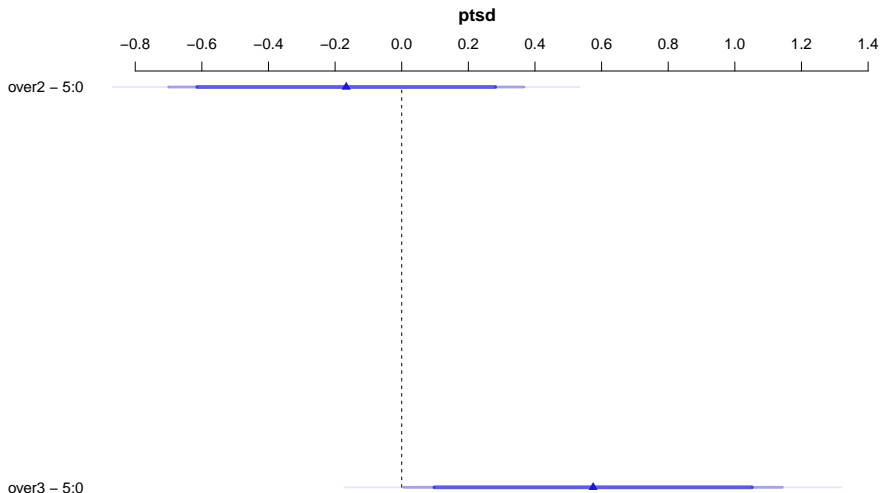
summary for mod_first fit by ols

```
summary(mod_first)
```

	Effects			Response : ptsd		
Factor	Low	High	Diff.	Effect	S.E.	Lower 0.95
over2	0	5	5	-0.16658	0.27195	-0.7103800
over3	0	5	5	0.57455	0.28970	-0.0047389
Upper 0.95						
0.37722						
1.15380						



Plot the summary to see effect sizes

```
plot(summary(mod_first))
```



What do the individual effects look like?

```
ggplot(Predict(mod_first))
```



Validate the summary statistics of an ols fit


```
set.seed(432010); validate(mod_first)
```

	index.orig	training	test	optimism
R-square	0.0633	0.0812	0.0309	0.0503
MSE	1.4893	1.4426	1.5408	-0.0982
g	0.3450	0.3607	0.3055	0.0552
Intercept	0.0000	0.0000	0.2893	-0.2893
Slope	1.0000	1.0000	0.8371	0.1629

	index.corrected	n
R-square	0.0130	40
MSE	1.5874	40
g	0.2899	40
Intercept	0.2893	40
Slope	0.8371	40


Build a nomogram for the ols fit

```
plot(nomogram(mod_first))
```



Is this model well-calibrated?

```
set.seed(432); plot(calibrate(mod_first))
```



n=64 Mean absolute error=0.224 Mean squared error=0.1048

0.9 Quantile of absolute error=0.64

Spending degrees of freedom wisely

- Suppose we have a data set with many possible predictors, and minimal theory or subject matter knowledge to guide us.
- We might want our final inferences to be as unbiased as possible. To accomplish this, we have to pay a penalty (in terms of degrees of freedom) for any “peeks” we make at the data in advance of fitting a model.
- So that rules out a lot of decision-making about non-linearity based on looking at the data, if our sample size isn't much larger than 15 times the number of predictors we're considering including in our model.
- In our case, we have $n = 64$ observations on 10 predictors.
- In addition, adding non-linearity to our model costs additional degrees of freedom.
- What can we do?

Spearman's ρ^2 plot: A smart first step?


Spearman's ρ^2 is an indicator (not a perfect one) of potential predictive punch, but doesn't give away the game.

- Idea: Perhaps we should focus our efforts re: non-linearity on predictors that score better on this measure.

```
spear_ptsd <- spearman2(ptsd ~ over2 + over3 + over5 + bond +  
                        posit + neg + contr + sup + cons + aff,  
                        data = maleptsd)
```

Spearman's ρ^2 Plot

```
plot(spear_ptsd)
```



Conclusions from Spearman ρ^2 Plot

- `neg` is the most attractive candidate for a non-linear term, as it packs the most potential predictive punch, so if it does turn out to need non-linear terms, our degrees of freedom will be well spent.
 - By no means is this suggesting that `neg` actually needs a non-linear term, or will show significant non-linearity. We'd have to fit a model with and without non-linearity in `neg` to know that.
 - Non-linearity will often take the form of a product term, a polynomial term, or a restricted cubic spline.
 - Since all of these predictors are quantitative, we'll think about polynomial or spline terms, soon.
- `over3`, also quantitative, has the next most potential predictive punch
- these are followed by `cons` and `aff`

Grim Reality

With 64 observations (63 df) we should be thinking about models with relatively tiny numbers of regression inputs.

- Non-linear terms (polynomials, splines) just add to the problem, as they need additional df to be estimated.

In this case, we might choose between

- including non-linearity in one (or maybe 2) variables (and that's it),
- or a linear model including maybe 3-4 predictors, tops

and even that would be tough to justify with this small sample size.

Contents of spear_ptsd

spear_ptsd


Spearman rho² Response variable:ptsd

	rho2	F	df1	df2	P	Adjusted rho2	n
over2	0.021	1.34	1	62	0.2522	0.005	64
over3	0.110	7.67	1	62	0.0074	0.096	64
over5	0.006	0.36	1	62	0.5527	-0.010	64
bond	0.004	0.22	1	62	0.6405	-0.013	64
posit	0.008	0.50	1	62	0.4825	-0.008	64
neg	0.136	9.73	1	62	0.0027	0.122	64
contr	0.001	0.03	1	62	0.8602	-0.016	64
sup	0.004	0.23	1	62	0.6357	-0.012	64
cons	0.052	3.40	1	62	0.0699	0.037	64
aff	0.052	3.39	1	62	0.0704	0.037	64

Actually Building Non-Linear Models


Predicting ptsd from over2

Linear and Loess Smooths of ptsd vs. over2



Linear Fit - does this work well?

```
plot(lm(ptsd ~ over2, data = maleptsd), which = 1)
```



Polynomial Regression

A polynomial in the variable x of degree D is a linear combination of the powers of x up to D .

For example:

- Linear: $y = \beta_0 + \beta_1x$
- Quadratic: $y = \beta_0 + \beta_1x + \beta_2x^2$
- Cubic: $y = \beta_0 + \beta_1x + \beta_2x^2 + \beta_3x^3$
- Quartic: $y = \beta_0 + \beta_1x + \beta_2x^2 + \beta_3x^3 + \beta_4x^4$
- Quintic: $y = \beta_0 + \beta_1x + \beta_2x^2 + \beta_3x^3 + \beta_4x^4 + \beta_5x^5$

Fitting such a model creates a **polynomial regression**.

Raw Quadratic Model for ptsd using over2

```
modA <- lm(ptsd ~ over2 + I(over2^2), data = maleptsd)
modA
```

Call:

```
lm(formula = ptsd ~ over2 + I(over2^2), data = maleptsd)
```

Coefficients:

(Intercept)	over2	I(over2^2)
1.23412	0.41128	-0.04213

$$ptsd = 1.234 + 0.411(over2) - 0.042(over2)^2$$

Summary of Quadratic Fit

```
> summary(modA)

Call:
lm(formula = ptsd ~ over2 + I(over2^2), data = maleptsd)


Residuals:
    Min       1Q   Median       3Q      Max
-2.0487 -1.2341  0.1503  0.9570  2.5945

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)   1.23412     0.24958   4.945 6.29e-06 ***
over2          0.41128     0.19271   2.134  0.0369 *
I(over2^2)    -0.04213     0.02014  -2.092  0.0407 *
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.246 on 61 degrees of freedom
Multiple R-squared:  0.0696,    Adjusted R-squared:  0.03909
F-statistic: 2.281 on 2 and 61 DF,  p-value: 0.1108
```

Plot Fitted Values of Quadratic Fit

Model A: Quadratic fit



Code for Previous Slide

```
modA_aug <- augment(modA, maleptsd)

ggplot(modA_aug, aes(x = over2, y = ptsd)) +
  geom_point() +
  geom_line(aes(x = over2, y = .fitted),
            col = "red") +
  labs(title = "Model A: Quadratic fit")
```

Another Way to fit the Identical Model

```
modA2 <- lm(ptsd ~ pol(over2, degree = 2, raw = TRUE),  
            data = maleptsd)
```

Coefficients:

(Intercept)	pol(over2, degree = 2, raw = TRUE)over2
1.23412	0.41128
pol(over2, degree = 2, raw = TRUE)over2^2	
-0.04213	

Do models give same fitted values?

```
temp <- fitted(modA2) - fitted(modA)  
sum(temp != 0)
```

```
[1] 0
```


Orthogonal Polynomials

Now, let's fit an orthogonal polynomial of degree 2 to predict ptsd using over2.


```
modB <- lm(ptsd ~ poly(over2, 2), data = maleptsd)
```

Looks very different ...


```
> modB  
  
Call:  
lm(formula = ptsd ~ poly(over2, 2), data = maleptsd)  
  
Coefficients:  
      (Intercept)  poly(over2, 2)1  poly(over2, 2)2  
          1.5925          0.5407          -2.6056
```

But it fits the same model, exactly!



Model A: Quadratic fit



Model B: Quadratic fit



Or, if you don't believe me...



Orthogonal Polynomial

An orthogonal polynomial sets up a model design matrix using the coding we've seen previously: `over2` and `over2^2` in our case, and then scales those columns so that each column is **orthogonal** to the previous ones.

- Two columns are orthogonal if their correlation is zero.
- This eliminates the collinearity (correlation between predictors) and lets our t tests tell us whether the addition of any particular polynomial term improves the fit of the model over the lower orders.

Would adding a cubic term help predict ptsd?

```
modC <- lm(ptsd ~ poly(over2, 3), data = maleptsd)
```

```
> summary(modC)
```

Call:

```
lm(formula = ptsd ~ poly(over2, 3), data = maleptsd)
```

Residuals:

Min	1Q	Median	3Q	Max
-1.9770	-1.1658	0.1784	0.9220	2.6628

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	1.5925	0.1567	10.164	1.15e-14	***
poly(over2, 3)1	0.5407	1.2534	0.431	0.6677	
poly(over2, 3)2	-2.6056	1.2534	-2.079	0.0419	*
poly(over2, 3)3	0.6363	1.2534	0.508	0.6135	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1


Residual standard error: 1.253 on 60 degrees of freedom

Multiple R-squared: 0.07357, Adjusted R-squared: 0.02725

F-statistic: 1.588 on 3 and 60 DF, p-value: 0.2016


Comparing Quadratic (red) and Cubic (blue) Models

Quadratic (red) vs. Cubic (blue) Polynomial Fits



What if we look instead at over3 as a predictor?

Linear and Loess Fits for ptsd vs. over3



What if we predict using over3?

```
modD1 <- lm(ptsd ~ over3, data = maleptsd)
modD2 <- lm(ptsd ~ poly(over3, degree = 2), data = maleptsd)
modD3 <- lm(ptsd ~ poly(over3, degree = 3), data = maleptsd)
```

```
> summary(modD1)$coef
```


	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.32814813	0.20649484	6.431871	2.047614e-08
over3	0.09723645	0.04999054	1.945097	5.630181e-02

```
> summary(modD3)$coef
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	1.592510	0.1460114	10.906746	7.216662e-16
poly(over3, degree = 3)1	2.419092	1.1680915	2.070979	4.267030e-02
poly(over3, degree = 3)2	-1.905737	1.1680915	-1.631496	1.080240e-01
poly(over3, degree = 3)3	3.225048	1.1680915	2.760955	7.635010e-03

Plotting the Fitted Models

Linear, Quadratic and Cubic Fits for ptsd using over3



Using Restricted Cubic Splines to Capture Non-Linearity

Splines

- A **linear spline** is a continuous function formed by connecting points (called **knots** of the spline) by line segments.
- A **restricted cubic spline** is a way to build highly complicated curves into a regression equation in a fairly easily structured way.
- A restricted cubic spline is a series of polynomial functions joined together at the knots.
 - Such a spline gives us a way to flexibly account for non-linearity without over-fitting the model.
 - Restricted cubic splines can fit many different types of non-linearities.
 - Specifying the number of knots is all you need to do in R to get a reasonable result from a restricted cubic spline.

The most common choices are 3, 4, or 5 knots.

- 3 Knots, 2 degrees of freedom, allows the curve to “bend” once.
- 4 Knots, 3 degrees of freedom, lets the curve “bend” twice.
- 5 Knots, 4 degrees of freedom, lets the curve “bend” three times.

Fitting Restricted Cubic Splines with `lm` and `rcs`

For most applications, three to five knots strike a nice balance between complicating the model needlessly and fitting data pleasingly. Let's consider a restricted cubic spline model for `ptsd` based on `over3` again, but now with:

- in `modE3`, 3 knots, and
- in `modE4`, 4 knots,





```
modE3 <- lm(ptsd ~ rcs(over3, 3), data = maleptsd)
modE4 <- lm(ptsd ~ rcs(over3, 4), data = maleptsd)
```

Summarizing the 4-knot model coefficients

Values of the estimates, and where are the knots located?

```
> round(summary(modE4)$coef,3)
              Estimate Std. Error t value Pr(>|t|)
(Intercept)      0.843      0.290    2.908   0.005
rcs(over3, 4)over3  0.953      0.432    2.209   0.031
rcs(over3, 4)over3' -8.914      5.982   -1.490   0.141
rcs(over3, 4)over3'' 13.480      9.635    1.399   0.167
>
> attributes(rcs(maleptsd$over3, 4))$parms
[1] 0.00 1.00 2.95 9.00
```

Plotting the spline models



Does the fit improve markedly from 3 to 4 knots?

In-sample comparison via ANOVA

```
anova(modE3, modE4)
```

Analysis of Variance Table

Model 1: `ptsd ~ rcs(over3, 3)`

Model 2: `ptsd ~ rcs(over3, 4)`

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	61	89.598				
2	60	87.573	1	2.0246	1.3871	0.2435

Does the fit improve markedly from 3 to 4 knots?

In-Sample comparisons of information criteria, etc.

```
glance(modE3) %>% select(r.squared, adj.r.squared, AIC, BIC)
```

```
# A tibble: 1 x 4
```

	r.squared	adj.r.squared	AIC	BIC
	<dbl>	<dbl>	<dbl>	<dbl>
1	0.119	0.0906	211.	220.


```
glance(modE4) %>% select(r.squared, adj.r.squared, AIC, BIC)
```

```
# A tibble: 1 x 4
```

	r.squared	adj.r.squared	AIC	BIC
	<dbl>	<dbl>	<dbl>	<dbl>
1	0.139	0.0963	212.	222.

Back to Spearman's ρ^2 Plot

```
plot(spear_ptsd)
```



Proposed New Model F

Fit a model to predict `ptsd` using:

- a 4-knot spline on `neg`
- a 3-knot spline on `over3`
- a linear term on `cons`
- a linear term on `aff`

Still more than we can reasonably do with 64 observations, but let's see how it looks.

Fit model F

```
modelF <- lm(ptsd ~ rcs(neg, 4) + rcs(over3, 3) +  
             cons + aff, data = maleptsd)
```

```
> round(summary(modelF)$coef,3)
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-0.425	0.749	-0.568	0.572
rcs(neg, 4)neg	0.066	0.060	1.095	0.278
rcs(neg, 4)neg'	-0.126	0.164	-0.768	0.446
rcs(neg, 4)neg''	0.492	0.537	0.916	0.363
rcs(over3, 3)over3	0.458	0.201	2.283	0.026
rcs(over3, 3)over3'	-2.125	0.943	-2.252	0.028
cons	-0.012	0.016	-0.724	0.472
aff	0.145	0.060	2.424	0.019

ANOVA for Model F

```
anova(modelF)
```

Analysis of Variance Table

Response: ptsd

	Df	Sum Sq	Mean Sq	F value	Pr(>F)	
rsc(neg, 4)	3	14.597	4.8657	3.7342	0.01617	*
rsc(over3, 3)	2	5.892	2.9460	2.2609	0.11369	
cons	1	0.636	0.6365	0.4885	0.48751	
aff	1	7.657	7.6566	5.8760	0.01860	*
Residuals	56	72.969	1.3030			

Signif. codes:

0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Remember that this ANOVA testing is sequential.

Is Model F better than Model E3?

```
anova(modelF, modE3)
```

Analysis of Variance Table

Model 1: $\text{ptsd} \sim \text{rcs}(\text{neg}, 4) + \text{rcs}(\text{over3}, 3) + \text{cons} + \text{aff}$

Model 2: $\text{ptsd} \sim \text{rcs}(\text{over3}, 3)$

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	56	72.969				
2	61	89.598	-5	-16.629	2.5524	0.03769 *

Signif. codes:

0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Limitations of `lm` for fitting complex linear models

We can certainly assess this big, complex model using `lm` in comparison to other models:

- with in-sample summary statistics like adjusted R^2 , AIC and BIC,
- we can assess its assumptions with residual plots, and
- we can also compare out-of-sample predictive quality through cross-validation,

But to really delve into the details of how well this complex model works, and to help plot what is actually being fit, we'll probably want to fit the model using `ols`, from the `rms` package.

Using `ols` to fit a complex linear model

Model F, fitted using `ols`

```
dd <- datadist(maleptsd)
options(datadist = "dd")

modF_ols <- ols(ptsd ~ rcs(neg, 4) + rcs(over3, 3) +
                cons + aff, data = maleptsd,
                x = TRUE, y = TRUE)
```


modF_ols results (slide 1 of 2)

```
> modF_ols
```

```
Linear Regression Model
```

```
ols(formula = ptsd ~ rcs(neg, 4) + rcs(over3, 3) + cons + aff,  
     data = maleptsd, x = TRUE, y = TRUE)
```

		Model Likelihood		Discrimination	
		Ratio Test		Indexes	
Obs	64	LR chi2	21.28	R2	0.283
sigma1.1415		d.f.	7	R2 adj	0.193
d.f.	56	Pr(> chi2)	0.0034	g	0.763

```
Residuals
```

	Min	1Q	Median	3Q	Max
	-2.06529	-0.81434	0.06745	0.81760	2.17200

modF_ols results (slide 2 of 2)

	Coef	S.E.	t	Pr(> t)
Intercept	-0.4255	0.7490	-0.57	0.5723
neg	0.0660	0.0603	1.10	0.2780
neg'	-0.1261	0.1641	-0.77	0.4456
neg''	0.4924	0.5373	0.92	0.3634
over3	0.4582	0.2007	2.28	0.0263
over3'	-2.1247	0.9433	-2.25	0.0282
cons	-0.0119	0.0164	-0.72	0.4722
aff	0.1450	0.0598	2.42	0.0186

Validation of Summary Statistics

```
set.seed(4322019); validate(modF_ols)
```

	index.orig	training	test	optimism
R-square	0.2829	0.3771	0.1484	0.2287
MSE	1.1401	0.9700	1.3539	-0.3838
g	0.7630	0.8665	0.6393	0.2272
Intercept	0.0000	0.0000	0.4268	-0.4268
Slope	1.0000	1.0000	0.7228	0.2772

	index.corrected	n
R-square	0.0542	40
MSE	1.5240	40
g	0.5358	40
Intercept	0.4268	40
Slope	0.7228	40

anova results for modF_ols

```
anova(modF_ols)
```

Analysis of Variance

Response: ptsd

Factor	d.f.	Partial SS	MS	F	P
neg	3	11.4062336	3.8020779	2.92	0.0420
Nonlinear	2	1.6536591	0.8268295	0.63	0.5339
over3	2	6.8378486	3.4189243	2.62	0.0814
Nonlinear	1	6.6106843	6.6106843	5.07	0.0282
cons	1	0.6826901	0.6826901	0.52	0.4722
aff	1	7.6565797	7.6565797	5.88	0.0186
TOTAL NONLINEAR	3	7.8079300	2.6026433	2.00	0.1248
REGRESSION	7	28.7821644	4.1117378	3.16	0.0070
ERROR	56	72.9688769	1.3030157		

summary results for modF_ols

```
summary(modF_ols)
```

Effects

Response : ptsd

Factor	Low	High	Diff.	Effect	S.E.	Lower	Upper
neg	11.175	30.425	19.25	0.44727	0.41704	-0.388160	0.95
over3	0.000	5.000	5.00	0.90059	0.43913	0.020902	
cons	45.750	55.000	9.25	-0.10997	0.15192	-0.414310	
aff	7.000	11.000	4.00	0.57998	0.23926	0.100680	

Upper 0.95

1.28270


1.78030

0.19437

1.05930


Plot of summary results for modF_ols

```
plot(summary(modF_ols))
```




Nomogram for modF_ols

```
plot(nomogram(modF_ols))
```




Seeing the impact of the modeling another way

```
ggplot(Predict(modF_ols))
```



Checking the model's calibration

```
set.seed(43220191); plot(calibrate(modF_ols))
```



n=64 Mean absolute error=0.153 Mean squared error=0.02813

0.9 Quantile of absolute error=0.236

Next Time

- The HERS data
- Fitting a more complex linear regression model
 - Dealing with categorical predictors
 - Dealing with interactions (another form of non-linearity)
 - Adding missing data into all of this, and running multiple imputation