432 Class 10 Slides

github.com/THOMASELOVE/2020-432

2020-02-18

Setup

```r
library(magrittr); library(janitor); library(here)

library(skimr)
library(naniar)
library(simputation)
library(broom)
library(rms) # note: also loads Hmisc
library(tidyverse)
```

Today's Data

Heart and Estrogen/Progestin Study (HERS)

- Clinical trial of hormone therapy for the prevention of recurrent heart attacks and deaths among 2763 post-menopausal women with existing coroneary heart disease (see Hulley et al 1998 and many subsequent references.)
- We're exluding the women in the trial with a diabetes diagnosis.

```
hers_raw <- read_csv(here("data/hersdata.csv")) %>%
    clean_names()

hers1 <- hers_raw %>%
    filter(diabetes == "no") %>%
    select(subject, ldl, ht, age, smoking, drinkany, sbp,
           physact, bmi, diabetes)
```

The Codebook (n = 2032)

Variable	Description
subject	subject code
HT	factor: hormone therapy or placebo
diabetes	yes or no (all are no in our sample)
ldl	LDL cholesterol in mg/dl
age	age in years
smoking	yes or no
drinkany	yes or no
sbp	systolic BP in mm Hg
physact	5-level factor, details next slide
bmi	body-mass index in kg/m^2

Goal Predict ldl using age, smoking, drinkany, sbp, physact and bmi, across both HT levels but restricted to women without diabetes.

The `physact` variable

```
hers1 %>% count(physact)
```

```
# A tibble: 5 x 2
  physact                n
  <chr>              <int>
1 about as active      674
2 much less active     107
3 much more active     252
4 somewhat less active 322
5 somewhat more active 677
```

Comparison is to activity levels for these women just before menopause.

Any missing data?

```
miss_var_summary(hers1)
```

```
# A tibble: 10 x 3
   variable n_miss pct_miss
   <chr>     <int>    <dbl>
 1 ldl           7   0.344
 2 drinkany      2   0.0984
 3 bmi           2   0.0984
 4 subject       0   0
 5 ht            0   0
 6 age           0   0
 7 smoking       0   0
 8 sbp           0   0
 9 physact       0   0
10 diabetes      0   0
```

Single Imputation for `drinkany`, `bmi` and `ldl`

Since `drinkany` is a factor, we have to do some extra work to impute.

```
set.seed(432092)

hers2 <- hers1 %>%
    mutate(drinkany_n =
                ifelse(drinkany == "yes", 1, 0)) %>%
    impute_pmm(drinkany_n ~ age + smoking) %>%
    mutate(drinkany =
                ifelse(drinkany_n == 1, "yes", "no")) %>%
    impute_rlm(bmi ~ age + smoking + sbp) %>%
    impute_rlm(ldl ~ age + smoking + sbp + bmi)
```

Now, check missingness...

```
miss_var_summary(hers2)
```

```
# A tibble: 11 x 3
   variable    n_miss pct_miss
   <chr>        <int>    <dbl>
 1 subject          0        0
 2 ldl              0        0
 3 ht               0        0
 4 age              0        0
 5 smoking          0        0
 6 drinkany         0        0
 7 sbp              0        0
 8 physact          0        0
 9 bmi              0        0
10 diabetes         0        0
11 drinkany_n       0        0
```

Multiple Imputation using `aregImpute` from `Hmisc`

Model to predict all missing values of any variables, using additive regression bootstrapping and predictive mean matching.

Steps are:

1. `aregImpute` draws a sample with replacement from the observations where the target variable is observed, not missing.
2. It then fits a flexible additive model to predict this target variable while finding the optimum transformation of it.
3. It then uses this fitted flexible model to predict the target variable in all of the original observations.
4. Finally, it imputes each missing value of the target variable with the observed value whose predicted transformed value is closest to the predicted transformed value of the missing value.

Fitting a Multiple Imputation Model

```
set.seed(4320132)
dd <- datadist(hers1)
options(datadist = "dd")
fit3 <- aregImpute(~ ldl + age + smoking + drinkany +
                     sbp + physact + bmi,
                 nk = c(0, 3:5), tlinear = FALSE,
                 data = hers1, B = 10, n.impute = 20)
```

Iteration 1 Iteration 2 Iteration 3 Iteration 4 Iteration 5 It

Multiple Imputation using `aregImpute` from `Hmisc`

`aregImpute` requires specifications of all variables, and several other details:

- `n.impute` = number of imputations, we'll run 20
- `nk` = number of knots to describe level of complexity, with our choice `nk = c(0, 3:5)` we'll fit both linear models and models with restricted cubic splines with 3, 4, and 5 knots
- `tlinear = FALSE` allows the target variable to have a non-linear transformation when `nk` is 3 or more
- `B = 10` specifies 10 bootstrap samples will be used
- `data` specifies the source of the variables

aregImpute **Imputation Results (1 of 4)**

```
fit3

Multiple Imputation using Bootstrap and PMM

aregImpute(formula = ~ldl + age + smoking + drinkany + sbp +
 physact + bmi, data = hers1, n.impute = 20, nk = c(0, 3:5),
 tlinear = FALSE, B = 10)

n: 2032     p: 7     Imputations: 20     nk: 0

Number of NAs:
 ldl  age  smoking drinkany  sbp  physact     bmi
   7    0        0        2    0        0       2
```

`fit3` Imputation Results (2 of 4)

```
R-squares for Predicting Non-Missing Values for Each
Variable Using Last Imputations of Predictors
     ldl drinkany      bmi
   0.041    0.014    0.109
```

`fit3` **Imputation Results (3 of 4)**

Resampling results for determining the complexity of imputation models

Variable being imputed: ldl

Bootstrap bias-corrrected summaries:

Statistic	nk = 0	nk = 3	nk = 4	nk = 5
R^2	0.0139	0.0149	0.00776	0.0124
mean absolute error	28.3594	42.9139	44.09937	39.8266
median abs. error	22.8301	35.5441	38.85302	32.6386

10-fold cross-validated:

Statistic	nk = 0	nk = 3	nk = 4	nk = 5
R^2	0.0214	0.0180	0.01517	0.0191
mean absolute error	145.7176	43.5007	45.02428	44.2456
median abs. error	141.4238	36.4102	38.88053	37.3141

`fit3` **Imputation Results (4 of 4)**

```
Variable being imputed: drinkany
                          nk=0   nk=3   nk=4    nk=5
Bootstrap  R^2           0.0163 0.0113 0.0102 0.00986
10-fold cv R^2           0.0205 0.0249 0.0163 0.01358
Bootstrap  mean |error|  0.4470 0.4568 0.4558 0.46624
10-fold cv mean |error|  0.4450 0.4454 0.4476 0.44676
Bootstrap median |error| 0.0000 0.0000 0.0000 0.00000
10-fold cv median |error| 0.0000 0.0500 0.1000 0.00000


Variable being imputed: bmi
                           nk=0    nk=3   nk=4   nk=5
Bootstrap  R^2            0.0845  0.0932 0.0946 0.0847
10-fold cv R^2           0.0864  0.0903 0.0968 0.0899
Bootstrap  mean |error|  3.7829  4.8119 4.9226 5.1775
10-fold cv mean |error|  27.6776 4.8359 4.9390 5.1136
Bootstrap median |error| 2.9955  3.9704 3.9371 4.2634
10-fold cv median |error| 27.0143 3.9894 3.9431 4.1876
```

A plot of the imputed values. . . (results)

A plot of the imputed values... (code)

```
par(mfrow = c(1,3))
plot(fit3)
par(mfrow = c(1,1))
```

- For ldl, we imputed most of the 7 missing subjects in most of the 20 imputation runs to values within a range of around 120 through 200, but occasionally, we imputed values that were substantially lower than 100.
- For drinkany we imputed about 70% no and 30% yes.
- For bmi, we imputed values ranging from about 23 to 27 in many cases, and up near 40 in other cases.
- This method never imputes a value for a variable that doesn't already exist in the data.

Kitchen Sink Model (Main Effects only)

```
mod_ks <- ols(ldl ~ age + smoking + drinkany + sbp +
              physact + bmi, data = hers2)
anova(mod_ks)
```

```
              Analysis of Variance          Response: ldl

 Factor      d.f.  Partial SS   MS           F      P
 age          1      9330.911   9330.911     6.93   0.0085
 smoking      1      8199.755   8199.755     6.09   0.0137
 drinkany     1      6444.424   6444.424     4.79   0.0288
 sbp          1      9274.287   9274.287     6.89   0.0087
 physact      4     10874.528   2718.632     2.02   0.0891
 bmi          1     15876.957  15876.957    11.80   0.0006
 REGRESSION   9     60077.708   6675.301     4.96  <.0001
 ERROR     2022   2721037.890   1345.716
```

# Spearman ρ^2 Plot

How should we prioritize the degrees of freedom we spend on non-linearity?

```
plot(spearman2(ldl ~ age + smoking + drinkany + sbp +
                physact + bmi, data = hers2))
```

Plot's on the next page.

- Note the use of the simple imputation hers2 data here. Why?

Spearman ρ^2 Plot Result

Spearman ρ^2 Response : ldl

Spending Degrees of Freedom

We're spending 9 degrees of freedom in our kitchen sink model. (We can verify this with anova or the plot.)

- Each quantitative main effect costs 1 df to estimate
- Each binary categorical variable also costs 1 df
- Multi-categorical variables with L levels cost L-1 df to estimate

Suppose we're willing to spend up to a total of **14** degrees of freedom (i.e. a combined 5 more on interaction terms and other ways to capture non-linearity.)

What should we choose?

What did we see in the Spearman ρ^2 Plot?

Group 1 (largest adjusted ρ^2)

- bmi, a quantitative predictor, is furthest to the right

Group 2 (next largest)

- smoking, a binary predictor, is next, followed closely by
- age, a quantitative predictor

Other predictors (rest of the group)

- sbp, quantitative
- drinkany, binary
- physact, multi-categorical (5 levels)

Impact of Adding Non-Linearity on Degrees of Freedom Spent

What happens when we add a non-linear term?

- Adding a polynomial of degree D costs D degrees of freedom.
    - So a polynomial of degree 2 (quadratic) costs 2 df, or 1 more than the main effect alone.
- Adding a restricted cubic spline with K knots costs K-1 df.
    - So adding a rcs with 4 knots uses 3 df, or 2 more than the main effect alone.
    - We restrict ourselves to considering splines with 3, 4, or 5 knots.
- Adding an interaction (product term) depends on the main effects of the predictors we are interacting
    - If the product term's predictors have df1 and df2 degrees of freedom, the product term adds df1 $\times$ df2 degrees of freedom to the main effects model.
    - An interaction of a binary and quantitative variable adds $1 \times 1 = 1$ additional degree of freedom to the main effects model.
    - When we use a quantitative variable in a spline and interaction, we'll do the interaction on the main effect, not the spline.

Model we'll fit with `ols`

Fitting a model to predict `ldl` using

- `bmi` with a restricted cubic spline, 5 knots
- `age` with a quadratic polynomial
- `sbp` as a linear term
- `drinkany` indicator
- `physact` factor
- `smoking` indicator and its interaction with the main effect of `bmi`

We can fit this to the data

- restricted to complete cases (hers1, effectively)
- after simple imputation (hers2)
- after our multiple imputation (fit3)

Fitting the model to the complete cases

```
d <- datadist(hers1)
options(datadist = "d")

m1 <- ols(ldl ~ rcs(bmi, 5) + pol(age, 2) + sbp +
              drinkany + physact + smoking +
              smoking %ia% bmi, data = hers1,
          x = TRUE, y = TRUE)
```

where %ia% identifies the linear interaction alone.

m1 results (slide 1 of 3)

```
> m1

Frequencies of Missing Values Due to Each Variable
 ldl   bmi   age   sbp drinkany physact smoking
   7     2     0     0        2       0       0

Linear Regression Model

ols(formula = ldl ~ rcs(bmi, 5) + pol(age, 2) + sbp +
    drinkany + physact + smoking + smoking %ia% bmi,
    data = hers1, x = TRUE, y = TRUE)

                 Model Likelihood      Discrimination
                   Ratio Test             Indexes
 Obs      2021   LR chi2      52.61   R2        0.026
 sigma36.7430   d.f.            14   R2 adj    0.019
 d.f.     2006   Pr(> chi2) 0.0000   g         6.629
```

m1 results (slide 2 of 3)

	Coef	S.E.	t	Pr(>|t|)
Intercept	121.6057	68.2000	1.78	0.0747
bmi	1.5687	1.0107	1.55	0.1208
bmi'	-8.6685	9.1577	-0.95	0.3440
bmi''	40.5712	37.4468	1.08	0.2787
bmi'''	-55.8872	44.5946	-1.25	0.2103
age	-0.5791	1.9657	-0.29	0.7683
age^2	0.0018	0.0149	0.12	0.9024
sbp	0.1221	0.0453	2.69	0.0072
drinkany=yes	-3.7427	1.6629	-2.25	0.0245
physact=much less active	-4.5660	3.8904	-1.17	0.2407
physact=much more active	-0.3291	2.7521	-0.12	0.9048
physact=somewhat less active	-0.0160	2.5270	-0.01	0.9950
physact=somewhat more active	3.7731	2.0293	1.86	0.0631
smoking=yes	-7.0832	12.0586	-0.59	0.5570
smoking=yes * bmi	0.4961	0.4391	1.13	0.2587

`m1` results (slide 3 of 3)

```
Residuals

     Min       1Q   Median       3Q      Max
-113.440  -24.519   -3.778   20.940  197.087
```

Fitting the model after simple imputation

```
dd <- datadist(hers2)
options(datadist = "dd")

m2 <- ols(ldl ~ rcs(bmi, 5) + pol(age, 2) + sbp +
              drinkany + physact + smoking +
              smoking %ia% bmi, data = hers2,
          x = TRUE, y = TRUE)
```

where, again, %ia% identifies the linear interaction alone.

`m2` results (slide 1 of 2)

```
> m2
Linear Regression Model

ols(formula = ldl ~ rcs(bmi, 5) + pol(age, 2) + sbp + drinkany +
    physact + smoking + smoking %ia% bmi, data = hers2, x = TRUE,
    y = TRUE)

                    Model Likelihood      Discrimination
                      Ratio Test             Indexes
 Obs      2032     LR chi2      53.14    R2        0.026
 sigma36.6503     d.f.            14    R2 adj    0.019
 d.f.     2017     Pr(> chi2) 0.0000    g         6.631

 Residuals

     Min       1Q    Median       3Q       Max
 -113.379  -24.326   -3.835   20.832   197.097
```

m2 results (slide 2 of 2)

```
                             Coef     S.E.      t    Pr(>|t|)
Intercept                 120.2662  67.6113   1.78   0.0754
bmi                         1.5508   1.0071   1.54   0.1237
bmi'                       -8.4486   9.0978  -0.93   0.3532
bmi''                      39.6413  37.1378   1.07   0.2859
bmi'''                    -54.8924  44.2677  -1.24   0.2151
age                        -0.5249   1.9490  -0.27   0.7877
age^2                       0.0014   0.0148   0.10   0.9233
sbp                         0.1209   0.0451   2.68   0.0074
drinkany=yes               -3.7023   1.6544  -2.24   0.0253
physact=much less active   -4.7408   3.8621  -1.23   0.2198
physact=much more active   -0.2635   2.7391  -0.10   0.9234
physact=somewhat less active 0.0130  2.5101   0.01   0.9959
physact=somewhat more active 3.8031  2.0193   1.88   0.0598
smoking=yes                -6.8961  12.0196  -0.57   0.5662
smoking=yes * bmi           0.4892   0.4375   1.12   0.2636
```

`anova(m2)` results

```
> anova(m2)
            Analysis of Variance          Response: ldl

Factor                                    d.f. Partial SS   MS         F    P
bmi  (Factor+Higher Order Factors)        5    2.758824e+04 5517.64861 4.11 0.0010
 All Interactions                         1    1.679813e+03 1679.81344 1.25 0.2636
 Nonlinear                                3    9.735452e+03 3245.15068 2.42 0.0647
age                                       2    9.175762e+03 4587.88077 3.42 0.0330
 Nonlinear                                1    1.244351e+01   12.44351 0.01 0.9233
sbp                                       1    9.657476e+03 9657.47569 7.19 0.0074
drinkany                                  1    6.726918e+03 6726.91809 5.01 0.0253
physact                                   4    9.709992e+03 2427.49791 1.81 0.1247
smoking  (Factor+Higher Order Factors)    2    1.085405e+04 5427.02463 4.04 0.0177
 All Interactions                         1    1.679813e+03 1679.81344 1.25 0.2636
smoking * bmi  (Factor+Higher Order Factors) 1 1.679813e+03 1679.81344 1.25 0.2636
TOTAL NONLINEAR                           4    9.738807e+03 2434.70175 1.81 0.1237
TOTAL NONLINEAR + INTERACTION             5    1.171134e+04 2342.26845 1.74 0.1214
REGRESSION                                14   7.178905e+04 5127.78931 3.82 <.0001
ERROR                                     2017 2.709327e+06 1343.24569
```

Validation of summary statistics

```
validate(m2)
```

```
          index.orig   training        test  optimism
R-square      0.0258     0.0322      0.0193    0.0129
MSE        1333.3300  1331.2658   1342.2160  -10.9502
g             6.6306     7.3074      5.9494    1.3581
Intercept     0.0000     0.0000     25.8305  -25.8305
Slope         1.0000     1.0000      0.8225    0.1775
          index.corrected  n
R-square           0.0129 40
MSE             1344.2801 40
g                  5.2726 40
Intercept         25.8305 40
Slope              0.8225 40
```

summary(m2) results

```
> summary(m2)
        Effects                     Response : ldl

Factor                                           Low   High    Diff.   Effect   S.E.   Lower 0.95 Upper 0.95
bmi                                              24.2  30.263  6.0625  5.1862  2.2217   0.82921    9.54330
age                                              62.0  72.000 10.0000 -3.3412  1.3450  -5.97890   -0.70357
sbp                                             120.0 145.000 25.0000  3.0218  1.1270   0.81165    5.23190
drinkany - yes:no                                 1.0   2.000     NA  -3.7023  1.6544  -6.94690   -0.45779
physact - about as active:somewhat more active    5.0   1.000     NA  -3.8031  2.0193  -7.76310    0.15695
physact - much less active:somewhat more active   5.0   2.000     NA  -8.5439  3.9035 -16.19900   -0.88862
physact - much more active:somewhat more active   5.0   3.000     NA  -4.0666  2.7125  -9.38630    1.25310
physact - somewhat less active:somewhat more active 5.0 4.000     NA  -3.7901  2.5633  -8.81720    1.23690
smoking - yes:no                                  1.0   2.000     NA   6.2635  2.4009   1.55500   10.97200

Adjusted to: bmi=26.9 smoking=no
```
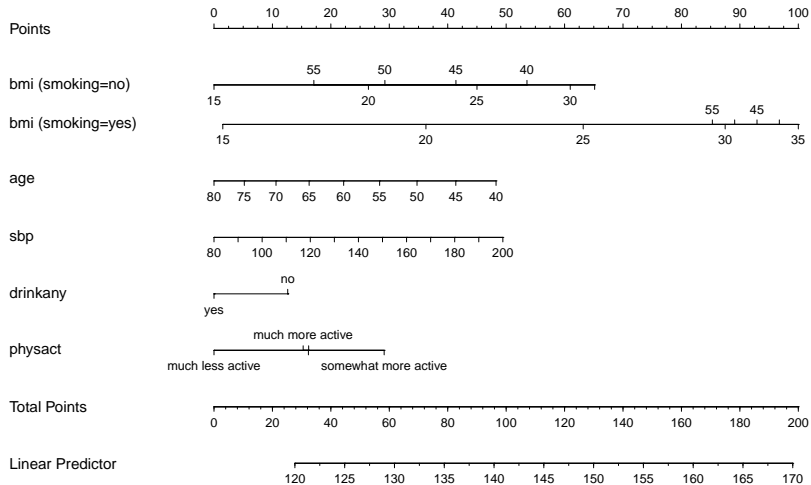
`plot(summary(m2))` results

Adjusted to:bmi=26.9 smoking=no

`plot(nomogram(m2))`

Making Predictions for an Individual

Suppose now that we want to use R to get a prediction for a new individual subject with bmi = 30, age = 50, smoking = yes and physact = about as active, drinkany= yes and sbp of 150.

```
predict(m2, expand.grid(bmi = 30, age = 50, smoking = "yes",
                        physact = "about as active",
                        drinkany = "yes", sbp = 150),
        conf.int = 0.95, conf.type = "individual")
```

```
$linear.predictors          $lower          $upper
         160.9399        88.48615       233.3936
```

Making Predictions for a Long-Run Mean

The other kind of prediction we might wish to make is for the mean of a series of subjects whose bmi = 30, age = 50, smoking = yes and physact = about as active, drinkany= yes and sbp of 150.

```
predict(m2, expand.grid(bmi = 30, age = 50, smoking = "yes",
                        physact = "about as active",
                        drinkany = "yes", sbp = 150),
        conf.int = 0.95, conf.type = "mean")
```

```
$linear.predictors         $lower       $upper
           160.9399      151.8119     170.0679
```

Of course, the confidence interval will always be narrower than the prediction interval given the same predictor values.

Influential Points?

```
which.influence(m2, cutoff = 0.4)
```

```
$Intercept
[1] 1135

$age
[1] 1135

$smoking
[1] 132

$`smoking * bmi`
[1] 132
```

Fitting the Model using Multiple Imputation

What do we have now?

- An imputation model fit3

```
fit3 <- aregImpute(~ ldl + age + smoking + drinkany + sbp +
          physact + bmi, nk = c(0, 3:5), tlinear = FALSE,
          data = hers1, B = 10, n.impute = 20, x = TRUE)
```

- A prediction model

```
m1 <- ols(ldl ~ rcs(bmi, 5) + pol(age, 2) + sbp +
            drinkany + physact + smoking + smoking %ia% bmi,
            x = TRUE, y = TRUE)
```

Now we put them together with the fit.mult.impute function...

Linear Regression & Imputation Model

```
m3imp <-
    fit.mult.impute(ldl ~ rcs(bmi, 5) + pol(age, 2) + sbp +
                        drinkany + physact + smoking +
                        smoking %ia% bmi,
                    fitter = ols, xtrans = fit3,
                    data = hers1)
```

```
Variance Inflation Factors Due to Imputation:

                    Intercept                              bmi
                         1.00                             1.00
                         bmi'                            bmi''
                         1.00                             1.01
                       bmi'''                              age
                         1.01                             1.00
                        age^2                              sbp
                         1.00                             1.01
```

m3imp results (1 of 2)

```
> m3imp
Linear Regression Model

 fit.mult.impute(formula = ldl ~ rcs(bmi, 5) + pol(age, 2) + sbp +
     drinkany + physact + smoking + smoking %ia% bmi, fitter = ols,
     xtrans = fit3, data = hers1)

                   Model Likelihood      Discrimination
                      Ratio Test            Indexes
 Obs       2032    LR chi2     53.30    R2        0.026
 sigma36.7128      d.f.           14    R2 adj    0.019
 d.f.      2017    Pr(> chi2) 0.0000    g         6.652

 Residuals

     Min      1Q  Median      3Q     Max
 -113.10  -24.46   -3.81   20.92  197.42
```

`m3imp` results (2 of 2)

```
                               Coef    S.E.     t    Pr(>|t|)
Intercept                   121.1499 67.7998  1.79  0.0741
bmi                           1.5445  1.0097  1.53  0.1263
bmi'                         -8.2945  9.1027 -0.91  0.3623
bmi''                        39.0890 37.3055  1.05  0.2949
bmi'''                      -54.2119 44.4779 -1.22  0.2230
age                          -0.5521  1.9547 -0.28  0.7776
age^2                         0.0016  0.0148  0.11  0.9119
sbp                           0.1216  0.0453  2.69  0.0073
drinkany=yes                 -3.7404  1.6625 -2.25  0.0246
physact=much less active     -4.7426  3.8692 -1.23  0.2204
physact=much more active     -0.2665  2.7455 -0.10  0.9227
physact=somewhat less active  0.0313  2.5214  0.01  0.9901
physact=somewhat more active  3.8060  2.0257  1.88  0.0604
smoking=yes                  -6.9198 12.0472 -0.57  0.5658
smoking=yes * bmi             0.4917  0.4388  1.12  0.2626
```

`anova(m3imp)`

```
> anova(m3imp)
              Analysis of Variance              Response: ldl

Factor                                          d.f. Partial SS    MS         F    P
bmi  (Factor+Higher Order Factors)               5   27514.6406 5502.9281 4.08 0.0011
 All Interactions                                1    1692.6044 1692.6044 1.26 0.2626
 Nonlinear                                       3    9741.6194 3247.2065 2.41 0.0653
age                                              2    9078.9851 4539.4926 3.37 0.0347
 Nonlinear                                       1      16.5032   16.5032 0.01 0.9119
sbp                                              1    9721.1667 9721.1667 7.21 0.0073
drinkany                                         1    6822.3861 6822.3861 5.06 0.0246
physact                                          4    9690.3632 2422.5908 1.80 0.1267
smoking  (Factor+Higher Order Factors)           2   10845.6127 5422.8063 4.02 0.0180
 All Interactions                                1    1692.6044 1692.6044 1.26 0.2626
smoking * bmi  (Factor+Higher Order Factors)     1    1692.6044 1692.6044 1.26 0.2626
TOTAL NONLINEAR                                  4    9747.0966 2436.7741 1.81 0.1246
TOTAL NONLINEAR + INTERACTION                   5   11717.3715 2343.4743 1.74 0.1225
REGRESSION                                      14   71571.1297 5112.2236 3.79 <.0001
ERROR                                         2017 2718570.0412 1347.8285
```

Evaluation via Partial R^2 and AIC (result)

432 Class 10 Slides

Evaluation via Partial R² and AIC (code)

```
par(mfrow = c(1,2))
plot(anova(m3imp), what="partial R2")
plot(anova(m3imp), what="aic")
par(mfrow = c(1,1))
```

summary(m3imp)

```
> summary(m3imp)
          Effects                        Response : ldl

Factor                                           Low    High    Diff.   Effect   S.E.   Lower 0.95 Upper 0.95
bmi                                              24.2   30.263  6.0625  5.2108  2.2283    0.84072   9.58080
age                                              62.0   72.000 10.0000 -3.3219  1.3498   -5.96910  -0.67463
sbp                                             120.0  145.000 25.0000  3.0394  1.1317    0.81989   5.25880
drinkany - yes:no                                 1.0    2.000      NA -3.7404  1.6625   -7.00080  -0.47996
physact - about as active:somewhat more active    5.0    1.000      NA -3.8060  2.0257   -7.77860   0.16663
physact - much less active:somewhat more active   5.0    2.000      NA -8.5486  3.9114  -16.21900  -0.87779
physact - much more active:somewhat more active   5.0    3.000      NA -4.0724  2.7198   -9.40640   1.26160
physact - somewhat less active:somewhat more active 5.0  4.000      NA -3.7746  2.5773   -8.82900   1.27980
smoking - yes:no                                  1.0    2.000      NA  6.3067  2.4196    1.56150  11.05200

Adjusted to: bmi=26.9 smoking=no
```
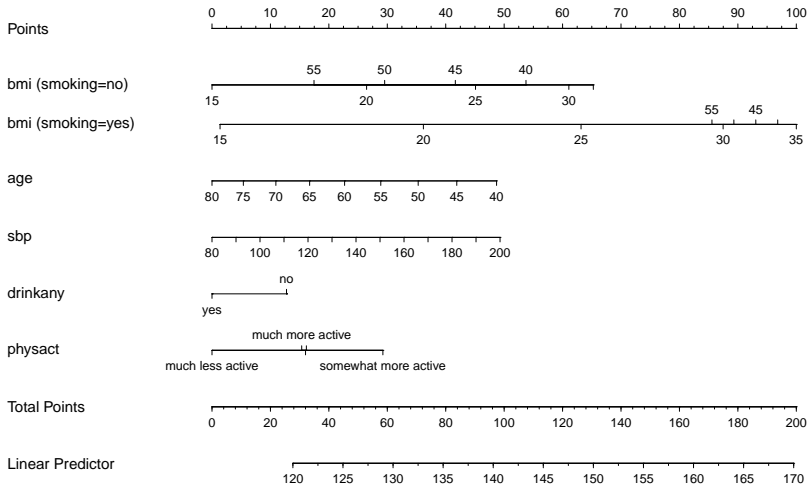
`plot(summary(m3imp))`

Adjusted to:bmi=26.9 smoking=no

plot(resid(m3imp) ~ fitted(m3imp))

`plot(nomogram(m3imp))`

Next Step

Can we do all of this for a logistic regression model?