

432 Class 19 Slides

github.com/THOMASELOVE/2020-432

2020-04-02

Today: Ordinal Logistic Regression and Using `rms`

Setup

```
library(knitr); library(janitor); library(magrittr)
library(caret)
library(rms)
library(nnet)
library(MASS)
library(broom)
library(tidyverse)
```

```
asbestos <- read_csv("data/asbestos.csv") %>% type.convert()
```

Asbestos Exposure in the U.S. Navy

These data describe 83 Navy workers, engaged in jobs involving potential asbestos exposure.

- The workers were either removing asbestos tile or asbestos insulation, and we might reasonably expect that those exposures would be different (with more exposure associated with insulation removal).
- The workers either worked with general ventilation (like a fan or naturally occurring wind) or negative pressure (where a pump with a High Efficiency Particulate Air filter is used to draw air (and fibers) from the work area.)
- The duration of a sampling period (in minutes) was recorded, and their asbestos exposure was measured and classified in three categories:
 - low exposure (< 0.05 fibers per cubic centimeter),
 - action level (between 0.05 and 0.1) and
 - above the legal limit (more than 0.1 fibers per cc).

Source Simonoff JS (2003) *Analyzing Categorical Data*. New York: Springer, Chapter 10.

Our Outcome and Modeling Task

We'll predict the ordinal Exposure variable, in an ordinal logistic regression model with a proportional odds assumption, using the three predictors

- Task (Insulation or Tile),
- Ventilation (General or Negative pressure) and
- Duration (in minutes).

Exposure is determined by taking air samples in a circle of diameter 2.5 feet around the worker's mouth and nose.

Summarizing the Asbestos Data

We'll make sure the Exposure factor is ordinal...

```
asbestos$Exposure <- factor(asbestos$Exposure, ordered=T)
summary(asbestos[,2:5])
```

Task	Ventilation	Duration
Insulation:46	General :34	Min. : 30.0
Tile :37	Negative pressure:49	1st Qu.: 85.0
		Median :138.0
		Mean :147.1
		3rd Qu.:212.5
		Max. :300.0

Exposure
(1) Low exposure :45
(2) Action level : 6
(3) Above legal limit:32

The Proportional-Odds Cumulative Logit Model

We'll use the `polr` function in the `MASS` library to fit our ordinal logistic regression.

- Clearly, Exposure group (3) Above legal limit, is worst, followed by group (2) Action level, and then group (1) Low exposure.
- We'll have two indicator variables (one for Task and one for Ventilation) and then one continuous variable (for Duration).
- The model will have two logit equations: one comparing group (1) to group (2) and one comparing group (2) to group (3), and three slopes, for a total of five free parameters.

Equations to be Fit

The equations to be fit are:

$$\log\left(\frac{\Pr(\text{Exposure} \leq 1)}{\Pr(\text{Exposure} > 1)}\right) = \beta_{0[1]} + \beta_1 \text{Task} + \beta_2 \text{Ventilation} + \beta_3 \text{Duration}$$

and

$$\log\left(\frac{\Pr(\text{Exposure} \leq 2)}{\Pr(\text{Exposure} > 2)}\right) = \beta_{0[2]} + \beta_1 \text{Task} + \beta_2 \text{Ventilation} + \beta_3 \text{Duration}$$

where the intercept term is the only piece that varies across the two equations.

- A positive coefficient β means that increasing the value of that predictor tends to *lower* the Exposure category, and thus the asbestos exposure.

Fitting the Model with the `polr` function in MASS

```
model.A <- polr(Exposure ~ Task + Ventilation + Duration,  
               data=asbestos)
```

Model Summary

```
> summary(model.A)
```

Re-fitting to get Hessian

Call:
polr(formula = Exposure ~ Task + Ventilation + Duration, data = asbestos)

Coefficients:

	Value	Std. Error	t value
TaskTile	-2.251333	0.644792	-3.4916
VentilationNegative pressure	-2.156979	0.567540	-3.8006
Duration	-0.000708	0.003799	-0.1864

Intercepts:

	value	Std. Error	t value
(1) Low exposure (2) Action level	-2.0575	0.6611	-3.1123
(2) Action level (3) Above legal limit	-1.5111	0.6344	-2.3820

Residual Deviance: 99.87952

AIC: 109.8795

Explaining the Model Summary

The first part of the output provides coefficient estimates for the three predictors.

	Value	Std. Error	t value
TaskTile	-2.251333	0.644792	-3.4916
VentilationNegative pressure	-2.156979	0.567540	-3.8006
Duration	-0.000708	0.003799	-0.1864

- The estimated slope for Task = Tile is -2.25. This means that Task = Tile provides less exposure than does the other Task (Insulation) so long as the other predictors are held constant.
- Typically, we would express this in terms of an odds ratio.

Odds Ratios and CI for Model A

```
exp(coef(model.A))
```

TaskTile	VentilationNegative pressure
0.1052589	0.1156740
Duration	
0.9992922	

```
exp(confint(model.A))
```

Waiting for profiling to be done...

Re-fitting to get Hessian

	2.5 %	97.5 %
TaskTile	0.02718379	0.3538549
VentilationNegative pressure	0.03641039	0.3427734
Duration	0.99187230	1.0069533

tidy for polr models exponentiates by default...

```
tidy(model.A, conf.int = TRUE)
```

term	estimate	std.error	statistic
(1) Low exposure (2) Action level	-2.057	0.661	-3.112
(2) Action level (3) Above legal limit	-1.511	0.634	-2.382
Duration	-0.001	0.004	-0.186
TaskTile	-2.251	0.645	-3.492
VentilationNegative pressure	-2.157	0.568	-3.801

term	conf.low	conf.high	coefficient_type
(1) Low exposure (2) Action level	NA	NA	zeta
(2) Action level (3) Above legal limit	NA	NA	zeta
Duration	-0.008	0.007	coefficient
TaskTile	-3.605	-1.039	coefficient
VentilationNegative pressure	-3.313	-1.071	coefficient

Assessing the Ventilation Coefficient

	Value	Std. Error	t value
TaskTile	-2.251333	0.644792	-3.4916
VentilationNegative pressure	-2.156979	0.567540	-3.8006
Duration	-0.000708	0.003799	-0.1864

Similarly, the estimated slope for Ventilation = Negative pressure (-2.16) means that Negative pressure provides less exposure than does General Ventilation. We see a relatively modest effect (near zero) associated with Duration.

Summary of Model A: Estimated Intercepts

Intercepts:

	Value	Std. Error	t value
(1) Low exposure (2) Action level	-2.0575	0.6611	-3.11
(2) Action level (3) Above legal limit	-1.5111	0.6344	-2.38

The first parameter (-2.06) is the estimated log odds of falling into category (1) low exposure versus all other categories, when all of the predictor variables (Task, Ventilation and Duration) are zero. So the first estimated logit equation is:

$$\log\left(\frac{Pr(Exposure \leq 1)}{Pr(Exposure > 1)}\right) =$$

$$-2.06 - 2.25[Task = Tile] - 2.16[Vent = NP] - 0.0007Duration$$

Summary of Model A: Estimated Intercepts

Intercepts:

	Value	Std. Error	t va
(1) Low exposure (2) Action level	-2.0575	0.6611	-3.1
(2) Action level (3) Above legal limit	-1.5111	0.6344	-2.3

The second parameter (-1.51) is the estimated log odds of category (1) or (2) vs. (3). The estimated logit equation is:

$$\log\left(\frac{Pr(Exposure \leq 2)}{Pr(Exposure > 2)}\right) =$$

$$-1.51 - 2.25[Task = Tile] - 2.16[Vent = NP] - 0.0007Duration$$

Comparing Model A to an “Intercept only” Model

```
model.1 <- polr(Exposure ~ 1, data=asbestos)
anova(model.1, model.A)
```

Likelihood ratio tests of ordinal regression models

Response: Exposure

	Model	Resid. df	Resid. Dev	Test
1	1	81	147.61971	
2	Task + Ventilation + Duration	78	99.87952	1 vs 2
	Df LR stat.	Pr(Chi)		
1				
2	3	47.74019	2.41857e-10	

Comparing Model A to Model without Duration

```
model.TV <- polr(Exposure ~ Task + Ventilation, data=asbestos)  
anova(model.A, model.TV)
```

Likelihood ratio tests of ordinal regression models

Response: Exposure

	Model	Resid. df	Resid. Dev	Test
1	Task + Ventilation	79	99.91421	
2	Task + Ventilation + Duration	78	99.87952	1 vs 2

	Df	LR stat.	Pr(Chi)
1			
2	1	0.03469471	0.8522368

Is a Task*Ventilation Interaction helpful?

```
model.TxV <- polr(Exposure ~ Task * Ventilation, data=asbestos)
anova(model.TV, model.TxV)
```

Likelihood ratio tests of ordinal regression models

Response: Exposure

	Model	Resid. df	Resid. Dev	Test	Df
1	Task + Ventilation	79	99.91421		
2	Task * Ventilation	78	99.64326	1 vs 2	1

LR stat. Pr(Chi)

1	
2	0.2709469 0.6026973

asbestos Likelihood Ratio Tests

Model	Elements	DF	Deviance	Test	<i>p</i>
1	Intercept	81	147.62	–	–
2	D	80	142.29	vs 1	0.021
3	T	80	115.36	vs 1	< 0.0001
4	V	80	115.45	vs 1	< 0.0001
5	T+V	79	99.91	vs 4	< 0.0001
6	T*V	78	99.64	vs 5	0.60
7	T+V+D	78	99.88	vs 5	0.85

- T = Task
- V = Ventilation
- D = Duration

In-Sample Predictions with our T+V model

```
model.TV <- polr(Exposure ~ Task + Ventilation,  
                 data=asbestos)  
  
asbestos <- asbestos %>%  
  mutate(TV_preds = predict(model.TV))  
  
asbestos %>% tabyl(TV_preds, Exposure) %>%  
  adorn_title() %>% kable()
```

	Exposure		
TV_preds	(1) Low exposure	(2) Action level	(3) Above legal
(1) Low exposure	42	3	10
(2) Action level	0	0	0
(3) Above legal limit	3	3	22

Accuracy of These Classifications?

```
asbestos %>% tabyl(TV_preds, Exposure) %>%  
  adorn_title() %>% kable()
```

	Exposure		
TV_preds	(1) Low exposure	(2) Action level	(3) Above legal
(1) Low exposure	42	3	10
(2) Action level	0	0	0
(3) Above legal limit	3	3	22

- Predicting Low exposure led to 42 right and 13 wrong.
- We never predicted Action Level
- Predicting Above Legal Limit led to 22 right and 6 wrong.

Total: 64 right, 19 wrong. Accuracy = $64/83 = 77.1\%$

5-fold cross-validation for polr model?

```
set.seed(2020432)
train.control <- trainControl(method = "cv", number = 5)
modTV_cv <- train(Exposure ~ Task + Ventilation,
                  data = asbestos, method = "polr",
                  trControl = train.control)
```

Results of 5-fold cross-validation modTV_cv

Ordered Logistic or Probit Regression

83 samples

2 predictor

3 classes: '(1) Low exposure', '(2) Action level', '(3) Above

No pre-processing

Resampling: Cross-Validated (5 fold)

Summary of sample sizes: 67, 66, 66, 67, 66

Resampling results across tuning parameters:

method	Accuracy	Kappa
cauchit	0.7477941	0.5069165
cloglog	0.7125000	0.4463752
logistic	0.7477941	0.5069165
loglog	0.7713235	0.5470383
probit	0.7477941	0.5069165

Which kappa is that?

Fleiss' kappa, or κ describes the extent to which the observed agreement between the predicted classifications and the actual classifications exceeds what would be expected if the predictions were made at random.

- Larger values of κ indicate better model performance ($\kappa = 0$ indicates very poor agreement between model and reality, κ near 1 indicates almost perfect agreement.)

Resampling results across tuning parameters:

method	Accuracy	Kappa
cauchit	0.7477941	0.5069165
cloglog	0.7125000	0.4463752
logistic	0.7477941	0.5069165
loglog	0.7713235	0.5470383
probit	0.7477941	0.5069165

Is the proportional odds assumption reasonable?

Alternative: fit a multinomial model?

```
mult_TV <- multinom(Exposure ~ Task + Ventilation,  
                    data = asbestos, trace = FALSE)
```

View the Multinomial Model?

```
mult_TV
```

Call:

```
multinom(formula = Exposure ~ Task + Ventilation, data = asbes,  
          trace = FALSE)
```

Coefficients:

	(Intercept)	TaskTile
(2) Action level	0.05268936	-1.160153
(3) Above legal limit	2.07821627	-2.699743
	VentilationNegative pressure	
(2) Action level		-2.316099
(3) Above legal limit		-2.496044

Residual Deviance: 98.08263

AIC: 110.0826

In-Sample Predictions with the multinomial T+V model

```
asbestos <- asbestos %>%  
  mutate(TVmult_preds = predict(mult_TV))  
  
asbestos %>% tabyl(TVmult_preds, Exposure) %>%  
  adorn_title() %>% kable()
```

	Exposure		
TVmult_preds	(1) Low exposure	(2) Action level	(3) Above legal
(1) Low exposure	42	3	10
(2) Action level	0	0	0
(3) Above legal limit	3	3	22

Compare Models with Likelihood Ratio Test?

```
(LL_multTV <- logLik(mult_TV)) # multinomial model: 6 df
```

```
'log Lik.' -49.04131 (df=6)
```

```
(LL_polrTV <- logLik(model_TV)) # polr model: 4 df
```

```
'log Lik.' -49.9571 (df=4)
```

```
(G = -2 * (LL_polrTV[1] - LL_multTV[1]))
```

```
[1] 1.831584
```

```
pchisq(G, 2, lower.tail = FALSE)
```

```
[1] 0.4001996
```

No statistically detectable difference in quality of fit ($p = 0.4$) between the proportional odds model and the more complex multinomial logistic regression model.

Using `rms` to fit ordinal logistic regression models

Proportional Odds Ordinal Logistic Regression with lrm

```
d <- datadist(asbestos)
options(datadist = "d")

model_TV_LRM <- lrm(Exposure ~ Task + Ventilation,
                    data = asbestos, x = TRUE, y = TRUE)

# note that Exposure must be an ordered factor
```

POLR results via lrm (slide 1)

```
model_TV_LRM
```

Logistic Regression Model

```
lrm(formula = Exposure ~ Task + Ventilation,  
     data = asbestos, x = TRUE, y = TRUE)
```

		Model Likelihood	
		Ratio Test	
Obs	83	LR chi2	47.71
(1) Low exposure	45	d.f.	2
(2) Action level	6	Pr(> chi2)	<0.0001
(3) Above legal limit	32		
max deriv	3e-10		

POLR results via lrm (slide 2)

```
lrm(formula = Exposure ~ Task + Ventilation + Duration,  
    data = asbestos, x = TRUE, y = TRUE)
```

Discrimination		Rank Discrim.	
Indexes		Indexes	
R2	0.526	C	0.854
g	2.064	Dxy	0.708
gr	7.877	gamma	0.839
gp	0.371	tau-a	0.396
Brier	0.127		

POLR results via lrm (slide 3)

```
lrm(formula = Exposure ~ Task + Ventilation + Duration,  
    data = asbestos, x = TRUE, y = TRUE)
```

	Coef	S.E.	Wald Z	Pr(> Z)
y>=(2) Action level	1.9713	0.4695	4.20	<0.0001
y>=(3) Above legal limit	1.4256	0.4348	3.28	0.0010
Task=Tile	-2.2868	0.6173	-3.70	0.0002
Ventilation=Negative pressure	-2.1596	0.5675	-3.81	0.0001

Ordinal Logistic Regression for T+V with orm

```
d <- datadist(asbestos)
options(datadist = "d")

model_TV_ORM <- orm(Exposure ~ Task + Ventilation,
                    data = asbestos, x = TRUE, y = TRUE)

# note that Exposure must be an ordered factor
```

Results for model_TV ORM fit with orm

(I'll neaten these up on the next two slides.)

```
model_TV_ORM
```

Logistic (Proportional Odds) Ordinal Regression Model

```
orm(formula = Exposure ~ Task + Ventilation, data = asbestos,  
    x = TRUE, y = TRUE)
```

		Model Likelihood Ratio Test	Discr In
Obs	83	LR chi2	47.71
(1) Low exposure	45	d.f.	2
(2) Action level	6	Pr(> chi2)	<0.0001
(3) Above legal limit	32	Score chi2	42.42
Distinct Y	3	Pr(> chi2)	<0.0001
Median Y	1		
max deriv	6e-05		

orm fit for T+V model (slide 1 of 2)

```
model_TV ORM
```

Logistic (Proportional Odds) Ordinal Regression Model

```
orm(formula = Exposure ~ Task + Ventilation,  
     data = asbestos, x = TRUE, y = TRUE)
```

		Model Likelihood	Ratio Test
Obs	83	LR chi2	47.71
(1) Low exposure	45	d.f.	2
(2) Action level	6	Pr(> chi2)	<0.0001
(3) Above legal limit	32	Score chi2	42.42
Distinct Y	3	Pr(> chi2)	<0.0001
Median Y	1		
max deriv	6e-05		

orm fit for T+V model (slide 2 of 2)

Logistic (Proportional Odds) Ordinal Regression Model

Discrimination Indexes

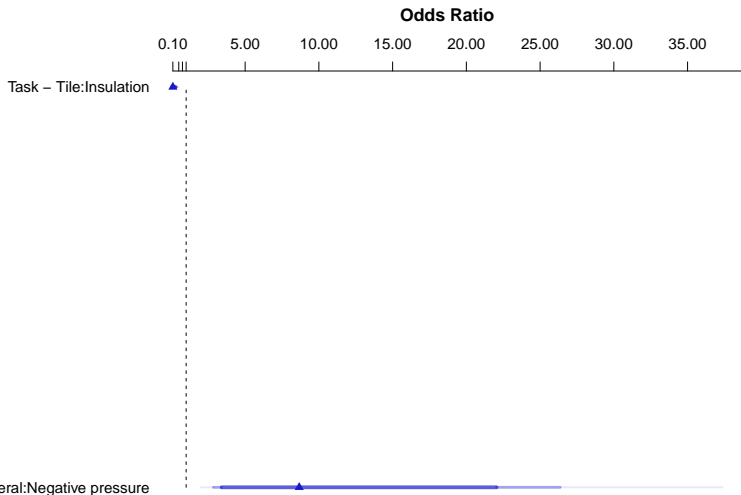
R2 0.526 rho 0.697

g 2.064 gr 7.877 |Pr(Y>=median)-0.5| 0.301

	Coef	S.E.	Wald Z	Pr(> Z)
y>=(2) Action level	1.9713	0.4695	4.20	<0.0001
y>=(3) Above legal limit	1.4256	0.4348	3.28	0.0010
Task=Tile	-2.2868	0.6173	-3.70	0.0002
Ventilation=Negative pressure	-2.1596	0.5675	-3.81	0.0001

Plot effects of the coefficients

```
plot(summary(model_TV_LRM))
```



rms::validate results from lrm

```
set.seed(432)
validate(model_TV_LRM)
```

	index.				index.	
	orig	training	test	optimism	corrected	n
Dxy	0.7077	0.7175	0.7082	0.0093	0.6984	40
R2	0.5260	0.5426	0.5183	0.0243	0.5017	40
Intercept	0.0000	0.0000	-0.0279	0.0279	-0.0279	40
Slope	1.0000	1.0000	0.9464	0.0536	0.9464	40
E _{max}	0.0000	0.0000	0.0169	0.0169	0.0169	40
D	0.5627	0.5944	0.5515	0.0429	0.5199	40
U	-0.0241	-0.0241	-0.4004	0.3763	-0.4004	40
Q	0.5868	0.6185	0.9519	-0.3335	0.9203	40
B	0.1270	0.1234	0.1319	-0.0086	0.1356	40
g	2.0639	2.1722	2.0250	0.1472	1.9167	40
gp	0.3709	0.3746	0.3691	0.0055	0.3654	40

rms::validate results from orm

```
set.seed(432)
validate(model_TV_ORM)
```

	index.orig	training	test	optimism	index.corrected
rho	0.6970	0.7113	0.6976	0.0138	0.6833
R2	0.5260	0.5470	0.5183	0.0287	0.4973
Slope	1.0000	1.0000	0.9354	0.0646	0.9354
g	2.0639	2.1963	2.0260	0.1702	1.8937
pdm	0.3010	0.3181	0.3042	0.0139	0.2871

n

rho	40
R2	40
Slope	40
g	40
pdm	40

Predictions (greater than or equal to)

```
head(predict(model_TV_LRM, type = "fitted"),3)
```

	y>=(2) Action level	y>=(3) Above legal limit
1	0.07762357	0.0464946
2	0.45306969	0.3243171
3	0.45306969	0.3243171

Predictions (individual)

```
head(predict(model_TV_LRM, type = "fitted.ind"),3)
```

	Exposure=(1) Low exposure	Exposure=(2) Action level
1	0.9223764	0.03112897
2	0.5469303	0.12875255
3	0.5469303	0.12875255

	Exposure=(3) Above legal limit
1	0.0464946
2	0.3243171
3	0.3243171

Nomogram?

First, we'll create the functions to estimate the probabilities of falling into groups 1, 2, and 3.

```
model_TV_LRM$coef
```

y>=(2) Action level	y>=(3) Above legal limit
1.971284	1.425557
Task=Tile Ventilation=Negative pressure	
-2.286807	-2.159559

So `plogis` by default uses the first intercept shown, and to get the machine to instead use the second one, we need:


```
fun3 <- function(x) plogis(x - model_TV_LRM$coef[2])
```

Plot the Nomogram

```
plot(nomogram(model_TV_LRM,  
  fun = list( 'Pr(y >= 2)' = plogis,  
              'Pr(y >= 3)' = fun3)))
```

Shown on next slide.


Points




Task




Ventilation




Total Points




Linear Predictor



$\Pr(y \geq 2)$



$\Pr(y \geq 3)$



Some Sources for Ordinal Logistic Regression

- A good source of information on fitting these models is <https://stats.idre.ucla.edu/r/dae/ordinal-logistic-regression/>
 - Another good source, that I leaned on heavily here, using a simple example, is <https://onlinecourses.science.psu.edu/stat504/node/177>.
 - Also helpful is <https://onlinecourses.science.psu.edu/stat504/node/178> which shows a more complex example nicely.